

Systematic Support of Parallel Bit Streams in LLVM

by

Meng Lin

B.Eng., University of Science and Technology of China, 2012

Thesis Submitted in Partial Fulfillment
of the Requirements for the Degree of

Master of Science

in the
School of Computing Science
Faculty of Applied Sciences

© Meng Lin 2014

SIMON FRASER UNIVERSITY

Fall 2014

All rights reserved.

However, in accordance with the *Copyright Act of Canada*, this work may be reproduced without authorization under the conditions for “Fair Dealing.” Therefore, limited reproduction of this work for the purposes of private study, research, criticism, review and news reporting is likely to be in accordance with the law, particularly if cited appropriately.

APPROVAL

Name: Meng Lin
Degree: Master of Science
Title of Thesis: Systematic Support of Parallel Bit Streams in LLVM

Examining Committee: Dr. Nick Sumner
Chair

Dr. Robert D. Cameron,
Professor, Computing Science,
Simon Fraser University
Senior Supervisor

Dr. Thomas C. Shermer,
Professor, Computing Science,
Supervisor

Dr. Fred Popowich,
Professor, Computing Science,
Supervisor

Dr. Arthur (Ted) Kirkpatrick,
Associate Director & Professor, Computing Science,
External Examiner

Date Approved: November 25th, 2014

Partial Copyright Licence



The author, whose copyright is declared on the title page of this work, has granted to Simon Fraser University the non-exclusive, royalty-free right to include a digital copy of this thesis, project or extended essay[s] and associated supplemental files ("Work") (title[s] below) in Summit, the Institutional Research Repository at SFU. SFU may also make copies of the Work for purposes of a scholarly or research nature; for users of the SFU Library; or in response to a request from another library, or educational institution, on SFU's own behalf or for one of its users. Distribution may be in any form.

The author has further agreed that SFU may keep more than one copy of the Work for purposes of back-up and security; and that SFU may, without changing the content, translate, if technically possible, the Work to any medium or format for the purpose of preserving the Work and facilitating the exercise of SFU's rights under this licence.

It is understood that copying, publication, or public performance of the Work for commercial purposes shall not be allowed without the author's written permission.

While granting the above uses to SFU, the author retains copyright ownership and moral rights in the Work, and may deal with the copyright in the Work in any way consistent with the terms of this licence, including the right to change the Work for subsequent purposes, including editing and publishing the Work in whole or in part, and licensing the content to other parties as the author may desire.

The author represents and warrants that he/she has the right to grant the rights contained in this licence and that the Work does not, to the best of the author's knowledge, infringe upon anyone's copyright. The author has obtained written copyright permission, where required, for the use of any third-party copyrighted material contained in the Work. The author represents and warrants that the Work is his/her own original work and that he/she has not previously assigned or relinquished the rights conferred in this licence.

Simon Fraser University Library
Burnaby, British Columbia, Canada

revised Fall 2013

Abstract

Parabix (parallel bit stream) technology uses the SIMD (single-instruction multiple-data) capabilities of commodity processors for high-performance text processing in applications such as Unicode transcoding, XML parsing and regular expression matching. LLVM is a widely-used compiler-infrastructure with a target-independent intermediate representation (IR) that includes notational support for SIMD operations on vectors of small integers. This thesis investigates whether it is possible to modify LLVM to incorporate all the SIMD processing requirements of Parabix both to increase the portability of applications and to create additional opportunities for optimization of those operations in the context of code generation. Our modifications to LLVM include redefining type legality and lowering for vectors of small elements as well as insertion of logic to recognize and properly handle Parabix-critical operations such as packing, merging and long stream addition. Experiments on the X86/SSE2 architecture show a speedup over the unmodified LLVM of about 300 times for some micro-benchmarks and demonstrate Parabix application performance slightly better than with its original SSE2 libraries. We also demonstrate performance scaling in switching from X86/SSE2 to X86/AVX2 without any change in source code.

To whomever whoever reads this!

“Don’t worry, Gromit. Everything’s under control!”
— *The Wrong Trousers*, AARDMAN ANIMATIONS, 1993

Acknowledgments

Here go all the people you want to thank.

Contents

Approval	ii
Partial Copyright License	iii
Abstract	iv
Dedication	v
Quotation	vi
Acknowledgments	vii
Contents	viii
List of Tables	x
List of Figures	xi
List of Programs	xii
Preface	xiii
1 Introduction	1
2 Background	4
2.1 SIMD and SWAR	4
2.1.1 Commercial SIMD Instructions Sets	4
2.2 Parabix Technology	5
2.2.1 IDISA Library	7
2.2.2 Critical Parabix Operations	9
2.3 LLVM Basics	10
2.4 LLVM Target-Independent Code Generator	11

2.5	Summary	13
3	Design Objectives	14
4	Vector of $i2^k$	21
4.1	Redefine Legality	22
4.2	In-place Lowering Strategy	23
4.2.1	Lowering for $vXi2$	23
4.2.2	Inductive Doubling Principle For $i4$ Vector	26
4.3	LLVM Vector Operation of $i2^k$	30
4.4	Long Stream Addition	32
5	Implementation	34
5.1	Standard Method For Custom Lowering	36
5.1.1	Custom Lowering Strategies	36
5.1.2	DAG Combiner	36
5.2	Templated Implementation	38
5.2.1	Code Generation For $i2$ Vector	38
5.2.2	Test Code And IR Library Generation	39
6	Performance Evaluation	41
6.1	Vector of $i2^k$ Performance	41
6.1.1	Methodology	41
6.1.2	Performance Against IDISA	42
6.1.3	Performance Against LLVM	44
6.2	Parabix Critical Operations	44
6.2.1	Ideal 3-Stage Transposition on the Intel Haswell	45
6.2.2	Long Stream Addition	46
7	Conclusion	50
	Bibliography	51

List of Tables

4.1	Supported operations and its semantics.	24
4.2	Legalize operations on $vXi1$ with iX equivalence.	24
4.3	Truth table of ADD on 2-bit integers and the minimized Boolean functions for C	25
4.4	Algorithm to lower $v32i4$ operations.	28
4.5	Rearranging index for BUILD VECTOR on $v64i2$	30
6.1	Hardware Configuration	43
6.2	Software Configuration	43
6.3	Performance against LLVM native support for $i2^k$ vectors	44
6.4	Performance comparison of XML Validator (xmlwf)	44
6.5	Performance comparison of UTF-8 UTF-16 Transcoder	45
6.6	Ideal 3-Stage Transposition with PEXT	45
6.7	Micro benchmarks for long stream addition against LLVM's original implementation.	46

List of Figures

2.1	Basis and Character Class Streams	6
2.2	ScanThru Using Bitstream Addition and Mask	7
2.3	MatchStar Using Bitstream Addition and Mask	7
2.4	Horizontal Operations in IDISA	8
2.5	Expansion Operations in IDISA	8
3.1	Implement <code>hsimd<32>::packh</code> with <code>shufflevector</code>	16
3.2	Implement <code>esimd<16>::mergeh</code> with <code>shufflevector</code>	17
4.1	Type legalize process for <i>v32i1</i> vector	22
4.2	Comparison between LLVM default legalize process and in-place lowering.	23
4.3	Addition of two <i>v32i4</i> vectors.	27
4.4	LLVM default type legalization of <i>v8i4</i> to <i>v8i8</i>	27
5.1	System overview: modified instruction selection process	35
5.2	Test system overview.	39
6.1	Test Performance with XOR	42
6.2	Reciprocal instruction throughput against IDISA library	47
6.3	Total CPU cycles against IDISA library	48
6.4	Vector of <i>i2</i> tested in a loop	49
6.5	Performance of <code>icgrep</code> with long stream addition	49

List of Programs

3.1	Implementation of <code>simd<8>::add</code> for X86 SSE2	14
3.2	Implementation of <code>simd<8>::add</code> for ARM NEON	15
3.3	Implementation of <code>simd<8>::add</code> with LLVM IR	15
3.4	Implementation of <code>simd<8>::eq</code> with LLVM IR	15
3.5	Implementation of <code>simd<8>::max</code> with LLVM IR	15
3.6	Shufflevector implementation of <code>packh</code>	17
3.7	Shufflevector implementation of <code>mergeh</code>	18
3.8	Clang-generated IR for <code>hsimd<8>::packh</code> from compiling the IDISA function	19
4.1	The function generated to lower ADD on <i>v64i2</i>	26
4.2	Inductive doubling principle on <i>v16i8</i> multiplication.	29
5.1	Implementation of <code>hsimd<2>::packh</code> with PEXT.	37
5.2	Templates for the IR Libray	40

Preface

Here go all the interesting reasons why you decided to write this thesis.

Chapter 1

Introduction

Nowadays Single Instruction Multiple Data (SIMD) instructions are broadly built in for most commodity processors. Compared with the traditional Single Instruction Single Data (SISD) instructions, SIMD provides an intra-register form of parallel computing by performing the same operation on many elements at the same time [17]. SIMD instructions are widely used for multimedia processing, digital signal processing or other compute-intensive applications [23, 25].

The recent method of parallel bit streams (Parabix) accelerates text processing using SIMD instructions, in applications such as UTF-8 to UTF-16 transcoding [12, 11], XML parsing [24, 14] and regular expression matching [26]. For these applications, byte streams of the input text characters are first transposed into 8 bit streams, one for each bit value of the character byte, and then loaded into SIMD registers so that 128 or 256 consecutive code units can be processed at once [15]. SIMD bitwise logic, shift operations, bit scans and other bit-based operations form the foundation of this programming model.

Although Parabix applications achieve substantial acceleration compared to sequential (SISD) equivalents, the Parabix tool chain needs to handle SIMD programming carefully. It is challenging for the following two major reasons:

1. SIMD instruction varies greatly among different instruction-set architectures (ISAs) which makes it hard to write portable SIMD programs. Some operations in Intel SSE may not exist in PowerPC AltiVec and vice versa. E.g. integer comparison intrinsic *pcmpgtq* in SSE4 does not have correspondence in PowerPC AltiVec.
2. Even within one specific SIMD instruction set, most operations are only available for some pre-chosen data sizes. This is referred as "sparse instruction set" in [10] and they gave a good example: in Intel SSE4, to shift-left a vector was implemented, but to shift-right was not. The 32-bit and 16-bit shift operations were available, but 64-bit shift was not [10].

The current Parabix tool chain uses the Inductive Doubling Instruction Set Architecture (IDISA) as an ideal computing model to overcome these two difficulties. Based on this model a library with the same name has been developed. It works well, but still has two shortcomings:

1. IDISA requires different header files for different architectures. Although it has a uniform API for portability, each header depends heavily on target-specific implementation details to provide the best implementation of each operation.
2. The IDISA generator [18] chooses the best implementation within the scope of a single function. This may be not the best when considering the context of this function. For example, when performing addition, we may know all the high bits of each field in a SIMD register is zero, making simplification possible.

This motivates us to find a better back end and currently, the most promising back end framework is LLVM. LLVM promises to enable out-sourcing of low-level and target-specific aspects of code generation [27, 21]. Switching to LLVM back end benefits Parabix tool chain for the following:

1. LLVM provides a target-independent intermediate representation (IR) with a vector of integer type system that matches IDISA requirement closely. Parabix operations can be expressed with IR and hence can be ported, in principle, to any platform that LLVM supports, including X86, ARM, PowerPC, MIPS, SPARC and many more.
2. LLVM provides inter-procedural, whole program analysis and optimization [22]. By built-in type system in the low level representation, LLVM keeps more static information to the back end and help optimize Parabix operation with meaningful context.
3. LLVM provides just-in-time compilation which allows runtime source generation. This is critical to some applications such as regular expression matching, which generates sequence of Parabix operations on the fly according to the input regular expression.

However, the native back end of LLVM has a number of gaps in its support for parallel bit streams. Important SIMD type of 2-bit and 4-bit field width are not available and 1-bit field width is supported slowly. Packing high bits on 16-bit field width which is one of the four key elements in the IDISA model and critical for transposition does not lower to proper machine code on X86. In this thesis, we extend LLVM to systematically support parallel bit streams and achieve high-performance code generation on the X86 target. We make the following contributions:

- We port the critical Parabix operations to the LLVM back end thus bringing all the benefit of LLVM discussed above into the Parabix technology.
- We redefine type legality in LLVM and extend LLVM type system with the inductive doubling principle so that vectors of small element type are properly supported.

- We insert logic in the LLVM back end to recognize and handle key Parabix operations. This allows efficient code generation while keeping the whole source code in target-independent IR.
- We add a dedicated LLVM intrinsic for the long stream addition and enable high-performance chained addition on the unbounded integer model which can be applied in broader applications.
- We evaluate the new LLVM back end with both micro benchmarks on single Parabix operation and application level profiles. We get the performance on X86 platform as good as the well-tuned IDISA library.

The remainder of this thesis is organized as follows. Chapter 2 provides a quick review of Parabix and LLVM. Chapter 3 shows the overall design goal; examples of the IR implementation are presented and compared with the IDISA library. Algorithms for machine code generation is discussed in Chapter 4 and the implementation details is in Chapter 5. Chapter 6 evaluates our work with both per-instruction benchmark and application level profile. Chapter 7 gives the conclusion.

Chapter 2

Background

2.1 SIMD and SWAR

SIMD is a parallel computing concept that performs the same instruction on different data to exploit data parallelism. Most of today's commodity processors supports SIMD within a register (SWAR). In this model, SIMD operations are applied within general-purpose or special registers that may be considered to be partitioned into fields. Operations on each field are independent from each other. This means for example, carry bits generated by addition could not pass to the next field.

The other important feature of the SWAR model is that the partition is not physical but rather a logical view of the register, so that different views are available on the same register. For a 128-bit SIMD register, a valid partition can be sixteen 8-bit fields as well as four 32-bit fields. There is no penalty from switching the logical view.

2.1.1 Commercial SIMD Instructions Sets

Some popular SIMD instructions sets are listed here:

- Intel MultiMedia eXtension (MMX). It defines eight 64-bit registers known as MM0 to MM7 which are aliases of the existing IA-32 Floating-Point Unit (FPU) stack registers. MMX only provides integer operations for early graphical applications thus is not a general purpose instruction set for SIMD programming [18].
- Intel Streaming SIMD Extensions (SSE) series. SSE extends the MMX instructions set and it introduces eight new independent 128-bit SIMD registers known as XMM0 to XMM7. Its successor SSE2 adds a rich set of integer instructions to the 128-bit XMM registers which makes it a useful SIMD programming model. AMD added support for SSE2 in its AMD64 architecture soon after the Intel released SSE2 thus in effect making SSE2 broadly available

across the desktop computers. Intel then released SSE3, SSSE3, SSE4 and AMD released SSE4a as the following SSE generations.

- Intel Advanced Vector Extensions (AVX). AVX extends the size of SIMD registers from 128 bits to 256 bits. It introduces 16 new registers YMM0 to YMM15. It fully supports SSE instructions and more importantly, shifts the two-operand operations towards the non-destructive three-operand form. Three-operand operation preserves the content in operand registers and could reduce the potential movement of data between registers. AVX supports a number of floating point operations on 256-bit registers, but does not support many of the integer operations that exist in SSE. Its successor AVX2 fills this gap and ensures the transition from SSE to AVX instructions with the same programming model. AVX2 is available on the Intel Haswell architecture. Its successor AVX512 has been announced to support 512-bit SIMD registers.
- ARM NEON. ARM as a popular mobile platform introduces its own SIMD extension named NEON in their Cortex-A series processors. It has thirty-two 64-bit registers (D0 to D31) as well as sixteen 128-bit registers (Q0 to Q15). In fact, $D_{2 \times i}$ and $D_{2 \times i + 1}$ are mapped to the same physical location of the register Q_i . Some operations like multiplication on the 64-bit D registers can return result in the 128-bit Q register [18]. NEON supports the field width of 8 bits, 16 bits, 32 bits and 64 bits integer operations as well as 32-bit floating point operations.

In this thesis, we use SSE2 as main ISA target with 128-bit registers because SSE2 is broadly available on both Intel and AMD CPUs. We use AVX2 as main ISA target with 256-bit SIMD registers because AVX lacks support of integer SIMD operations.

2.2 Parabix Technology

Parabix technology is a programming framework for high-performance text processing that can utilize both SIMD and multi-core parallel processing facility. It is built on top of the parallel bit streams concept. Byte-oriented input stream is first transposed into 8 bit streams with each stream corresponds to one bit location in the byte stream. For encodings that requires more than one byte, more bit streams can be introduced: one bit of the input code unit for one bit stream. Figure 2.1 gives an example of the transposition, B_0 to B_7 are the bit streams of the ASCII encoded input data. Zero bits are marked as periods (.) for clarity.

After the transposition, the character class bit streams are generated using bitwise logic, e.g. [a], [z9] and [0-9] in the figure. With SIMD operations on the 128-bit register, 128 input code unit can be classified at the same time. Parabix defines a set of primitives on the arbitrary length bit stream, called the *Pablo Language* which is usually applied on the character class bit streams to generate a number of *Marker Streams*. Marker Streams mark meaningful locations such as where

input data	a453z--b3z--az--a12949z--ca22z7--
B_7
B_6	1...1...1...11...1....1...11...1...
B_5	111111111111111111111111111111111111
B_4	.1111...11...1...11111...1111..
B_3	...111...111...111...1...1111...1.11
B_2	.11...11...11...11...1...1...11....111
B_1	...11...111...1...1...1...1...1.1111..
B_0	1.11.11.1.111.1111.1.1.1111...111
[a]	1.....1...1.....1.....
[z9]	...1...1...1...1...1.11....1...
[0-9]	.111....1.....11111....11.1..

Figure 2.1: Basis and Character Class Streams. Cited from [26].

a tag starts and ends in the XML document, matching positions of a partial regular expression. A simple counting or scanning through the marker streams are usually the final step in the Parabix technology.

Some useful Pablo primitives are listed as the following [13]:

1. Bitwise logic: AND, OR, XOR and NOT on arbitrary length bit stream.
2. Advance: shift forward the whole bit stream for 1 bit. In a little ending system, shift forward is to shift left because the bytes that comes first in the input stream resides in the lower memory address.
3. ScanThru: $s(M, C)$ denotes the operation of scanning from the marker stream M as the initial positions through the spans of ones in the stream C . Figure 2.2 shows an example of it.

$$s(M, C) = (M + C) \wedge \neg C$$

One example of ScanThru is in XML well-formedness checking, to check if a `<tag>` is written in correct syntax, M marks all the start positions of tags e.g. the next position after the opening angle bracket (`<`) and C is the marker stream for all legal tag content. $s(M, C)$ should mark all the positions of the closing angle bracket (`>`) which close tags. Say M_0 denotes the character class of `>`, then if $s(M, C) \wedge \neg M_0$ is not all zero, some tag is not closed properly with the `>` symbol. Note that all the tags in the stream are checked simultaneously in parallel in the unbounded bit streams model.

4. MatchStar: $m(M, C)$ returns all positions that can be reached by scanning from the initial positions marked in M along the spans of ones in the stream C for zero or more steps.

MatchStar gets its name from the star operator (*) in the regular expression. It also has important application in long stream addition. Figure 2.3 shows an example of it.

input data	----173942---654----1----49731----321--
M_01.....1.....1.....1.....
$D = [0-9]$111111...111....1....11111...111..
$M_0 + D$1.....1....1...11...1...111..
$M_1 = (M_0 + D) \wedge \neg D$1.....1....1.....1.....

Figure 2.2: ScanThru Using Bitstream Addition and Mask. Cited from [13] and slightly modified.

input data	a453z--b3z--az--a12949z--ca22z7--
M_1	.1.....1...1.....1.....
$C = [0-9]$.111....1.....11111....11.1..
$T_0 = M_1 \wedge C$.1.....1.....1.....
$T_1 = T_0 + C$1...1.....1.....11..
$T_2 = T_1 \oplus C$.1111.....111111...111...
$M_2 = T_2 \vee M_1$.1111.....1...111111...111...

Figure 2.3: MatchStar primitive, where $M_2 = \text{MatchStar}(M_1, C)$. Cited from [26].

The Pablo language is defined over unbounded bit streams which of course need to be translated into a block-at-a-time processing for real applications [26]. The Pablo compiler is used here for the translation and it takes care of the carry bits across block boundaries with a carry queue. A block-at-a-time C++ code is generated as a result.

2.2.1 IDISA Library

To actually execute the C++ code, a set of runtime library is necessary. Cameron proposed the Inductive Doubling Instructions Set Architecture in [15] and claimed significant instruction count reduction in core parallel bit stream algorithm. As he wrote, "inductive doubling refers to a general property of certain kinds of algorithm that systematically double the values of field widths or other data attributes with each iteration." [15]. There are four key elements of this architecture:

- A core set of binary functions on SIMD registers, for all field width equals to 2^k . To work with parallel bit streams, the operation ADD, SUB, SHL, SRL and ROTL (rotate left) comprise the set.

- A set of *half-operand modifiers* that make possible the inductive processing of field width $2W$ in terms of combinations of field width W . These modifiers select either the lower half of the field or the higher half.
- Packing operations that compress two vectors of field width W into one vector of field width $W/2$. E.g. collecting all the higher half bits of fields from two vectors into one.
- Merging operations that produce one vector of field width W with two vectors of field width $W/2$.

A C++ library is then developed after this model and it is called the IDISA library. To be clear, in the following sessions the abstract architecture is called the IDISA model to distinguish from the IDISA library. An interesting fact about the IDISA library is that it is actually generated automatically from a pool of strategies to avoid duplicated human work among different targets. When targeting at a new platform, the natively supported SIMD instructions need to be mapped to proper operations in the IDISA model. This is sparse in the sense that many other operations defined in the model are still not available. The IDISA generator could fill the gaps with a pool of strategies which basically tells how to implement instruction C given instruction A and B are available. Multiple strategies for the same operation may exist and the generator chooses based on the least instruction count mechanism [18].




Figure 2.4: The logic of IDISA Horizontal Operations, cited from [18].




Figure 2.5: The logic of IDISA Expansion Operations, cited from [18].

The IDISA library divides SIMD operations in the following categories [2]:

- Vertical Operations (Template Class `simd<w>`): Most common SIMD operations between two registers where `w` is the field width. E.g. `simd<8>::add(A, B)` aligns registers A and B vertically and adds up the aligned 8-bit fields. Different fields are independent of each other.
- Horizontal Operations (Template Class `hsimd<w>`): Operations like packing align the two operands horizontally, extract a portion of the bits in operands and concatenate into one full SIMD register. (Figure 2.4)
- Expansion Operations (Template Class `esimd<w>`): Operations that double the width of fields like merging which takes the higher 64 bits of the two operands (A and B), concatenates the first field from A and the first field from B to get a new field with the width doubled. Do the same to the following fields until a full SIMD register C is generated. (Figure 2.5)
- Field Movement Operations (Template Class `mvmd<w>`): Operations that copy and move the entire fields. The content of these fields should not change.
- Full Register Operations (Template Class `bitblock`): Operations that works with the contents of SIMD registers as undivided bit blocks.

The IDISA library claims to have better performance compared to the hand-written libraries and it is the main competitor of our LLVM back end.

2.2.2 Critical Parabix Operations

There are at least four critical Parabix operations that can be the performance bottleneck and need special attention:

- Transposition. The first step of every Parabix application and it can be the primary overhead of some Parabix application. There are two major algorithms: the ideal three-stage implementation and the byte-pack implementation. The byte-pack implementation utilize packing on 16 bit field width which is widely available on commodity processors. The ideal three-stage implementation is less available but it is proved optimal in instruction count in the IDISA model. Details of these two algorithms can be found in [15]. We implement both of the algorithm in IR and compare their performance in Chapter 6.
- Inverse transposition. For some applications like the UTF-8 to UTF-16 transcoding, parallel bit streams needs to be modified and translated back into byte streams, thus an inverse transposition is needed. As the inverse operation to the transposition, there are also two algorithms available which mirror the transposition algorithms. A detailed discussion can be found in [12].
- Long stream addition. Pablo compiler deal with addition between unbounded bit streams using chained long stream additions, which adds two numbers as wide as the SIMD register

with a carry-in bit and generates a carry-out bit. The naive approach chains 64-bit additions together to emulate 128-bit, 256-bit or 512-bit additions. Its time complexity grows linearly with the SIMD register size. A better algorithm is proposed in [26] which could add up to 4096 bits wide integers in constant time. We discuss this algorithm further in Chapter 4.

- Parallel deletion. This operation deletes bits from one bit stream at the positions identified by a deletion marker. Three different inductive doubling algorithms can be used for it. Refer to [12] for further detail.

These operations takes most of the run time in the application level profile, so a slight improvement could benefit the application greatly. We study the LLVM support of the first three operations in this work and leave the parallel deletion to be future work.

2.3 LLVM Basics

The *Low Level Virtual Machine (LLVM)* is an open-source, well developed compiler tool that is dedicated to the compiler writers. It originates in 2000 with Lattner's Master thesis [21] and is gaining popularity ever after. Today it is developed into a high-performance static compiler back end with just-in-time compilers and life-long program analysis and optimization, which means program analysis and optimization in compile time, link time and run time [27, 22]. It supports a variety of targets from Intel X86, PowerPC to the ARM mobile platform and hides the low-level target-specific issues for the compiler writers.

LLVM defines an intermediate representation (IR) as its virtual instruction set and IR is used as not only the input code to the LLVM tool chain but also the internal representation for analysis and optimization passes. This enables the programmer to use LLVM as a pipeline and to inspect output from each step. For example, with the C/C++ compiler of LLVM, Clang, source code in C/C++ is compiled into IR with different level of front end optimization. IR files are then linked together with link-time optimizations such as function inline. The result IR is put through a good number of LLVM passes. LLVM optimizations are organized into passes. There are three types of passes [5]: analysis passes that collect information for other passes, transform passes that change the program and utility passes that provide general helper functionality for other passes. Each of the pass consumes LLVM IR as input and generate IR output. The mutation it makes to the file can be easily checked by running the pass alone.

Although the IR is low-level, it preserves high-level static information through the strong type system and static single assignment (SSA) form. SSA form guarantees only one assignment to every variable. According to [16], SSA helps calculate the high-level data flow. The main design goal of the IR is to be low-level enough so that most programming language can target to it while maintaining the most high-level information to make aggressive back end optimization possible [27].

LLVM IR defines instructions and intrinsic functions. There are terminator instructions which produce control flow (return, branch, etc.), binary instructions (add, subtract, etc.), bitwise binary instructions (shift left/right, logic operations, etc.), memory instructions (load, store, etc.) and other instructions. Intrinsic functions are extension of IR instructions. Their name must all start with an "llvm." prefix [1]. Example intrinsic functions are standard C library intrinsics (memcpy, sqrt, sin, floor, etc.), bit manipulation intrinsics (population count, byte swap, etc.), debugger intrinsics, exception handling intrinsics and so on. They can be general operations for all the platform as well as target-specific; e.g. `llvm.x86.sse2.psrlw` corresponds to SSE2 native instruction `psrlw`. We use none of these target-specific intrinsics in our work to achieve portability.

After optimization passes, the IR code are processed through the target-independent code generator and the machine code (MC) layer to become the native machine code. We describe the code generation process in detail in the next section as it is the major piece of logic we extend for parallel bit streams.

2.4 LLVM Target-Independent Code Generator

The first stage for code generation is Instruction Selection, which translates LLVM code into the target-specific machine instructions. After that, there are machine level optimization like live-interval analysis and register allocation. Instruction Selection is done by the following steps [4] (we describe each step in the following text):

- Initial SelectionDAG Construction: generate SelectionDAG from LLVM IR.
- DAG Combine 1
- Type Legalization Phase
- Post Type Legalization DAG Combine
- Operation Legalization Phase
- DAG Combine 2
- Instruction Select Phase
- Scheduling and Formation Phase

LLVM internally constructs a graph view of the input code called SelectionDAG where DAG is short for directed acyclic graph. Each node in the DAG represents an operation with an opcode, a number of operands and a number of return values. If the DAG node A uses the return value of another DAG node B, there will be an edge from B to A. The SelectionDAG enables a large variety

of very-low-level optimization. And it also benefits the instruction scheduling process by recording the instruction dependency in the graph.

There are DAG combine passes after the initial construction and each legalization phase[4]. We will explain "legality" in the next section. DAG combine passes clean up SelectionDAG with both general and machine-dependent strategies, making the work easier for initial constructor and legalizers: they can focus on generating accurate SelectionDAG, good and legal operations with no worries of the messy output.

Instruction Select Phase is the bulk of target-specific logic that translates a legal SelectionDAG into a new DAG of target code with pattern matching facility. For example, a node of floating point addition followed by a floating point multiplication could be merged into one FMADDS node on the target that supports floating point multiply-and-add (FMA) operations [4].

The Scheduling and Formation Phase assigns an order to each target instruction following the target's constraints. After that, a list of machine instructions are generated and the SelectionDAG is no longer needed.

Now we look at how LLVM deals with SIMD instructions. SIMD data are grouped into vectors and LLVM uses the notion $\langle N \times iX \rangle$ to represent a vector of N elements, where each of the element is an integer of X bits [1, 10]. $\langle N \times iX \rangle$ is also denoted as $vNiX$ as $vNiX$ is the internal type name used in the LLVM source code; e.g. $\langle 4 \times i32 \rangle$ is the same with $v4i32$. Various of operations can be applied on vectors and the LLVM back end knows how to lower them into proper machine instructions.

In LLVM IR, programmers can write any kind of vectors, even $v1024i3$. Those vectors may not be supported by the target machine. LLVM has the notion of "legal" vs. "illegal" types. Legality is an target-specific concept. A DAG node is legal if it only uses the supported operation on the supported types. Unsupported types are illegal types for the target. For example, $i1$ is not supported in X86, it is illegal together with all the operations that take $i1$ operands or return $i1$ values. Addition on $v16i8$ is legal for X86 SSE2 but multiplication on $v16i8$ is not since there is no native support of it. The type $v16i8$ is considered to be legal in this case. LLVM code generator has all the target details. It uses type legalization and operation legalization phases to turn illegal type or DAG into legal[4].

Type legalization phase has three ways to legalize vector types[10]: *Scalarization*, *Vector Widening* and *Vector Element Promotion*.

- **Scalarization** splits the vector into multiple scalars. It is often used for $v1iX$ as the edge case when LLVM is trying to legalize by splitting the incoming vector into sub vectors.
- **Vector Widening** adds dummy elements to make the vector fit the right register size. It does not change the type of the elements, e.g. $v4i8$ to $v16i8$.
- **Vector Element Promotion** preserves the number of elements, but promote the element type to a wider size, e.g. $v4i8$ to $v4i32$.

After type legalization, we may still have illegal DAG node; thus we need legalize operations phase. There are three strategies in this phase:

- **Expansion:** Use another sequence of operations to emulate the operation. Expansion strategy is often general in the sense that it may use slow operations such as memory load and store, but it generates native code with correct outcome.
- **Promotion:** Promote the operand type to a larger type that the operation supports.
- **Custom:** Write a target-specific code to implement the legalization. Similar to Expansion, but with a specific target in mind.

No illegal type should be introduced in the operation legalization phase which puts a limitation on the machine-independent legalize strategies: *i8* is the minimum integer type on X86 and programmer needs to extend every integer less than 8 bits to *i8* before returning it to the DAG. On the other hand DAG combine is different, you can choose the combine timing on your own. If you choose to combine before type legalization phase, you can freely introduce illegal types into your combined results.

The current legalization mechanism of LLVM is not sufficient to handle Parabix code efficiently. We propose new strategies and redefine legality in Chapter 4.

2.5 Summary

In this chapter we reviewed the Parabix technology which is a parallel text processing model and IDISA library which is a C++ library for SIMD programming. We also provided LLVM basics and described its target-independent code generator. IDISA library has to maintain target-specific header files and is hard to further optimize outside the function scope. So we propose to replace IDISA library with a LLVM back end in the next chapter.

Chapter 3

Design Objectives

In this chapter we discuss about our overall goal for using LLVM as a new Parabix back end. First, we show that the IDISA library could be replaced by a pure target-independent IR library.

To start, let us look at one IDISA vertical operation: `simd<8>::add`. IDISA library implements this function with the compiler intrinsic that directly translates into the assembly code, so different header files have to be maintained for different instruction sets such as Program 3.1 and Program 3.2. However, with the LLVM IR, we can implement it as Program 3.3; no low level detail is specified here.

Most of the IDISA vertical operations can be expressed with a few lines of IR code. A bit more examples are listed here:

- Vector addition, subtraction, multiplication and shifting. There are IR instructions that correspond one-to-one with them.
- Integer comparison such as equality, greater than and unsigned less than. In IR, there is one instruction called 'icmp' which does the comparison. The only difference is that for vector type `<N x iX>`, the comparison result of 'icmp' is in type `<N x i1>` while IDISA requires it to be in type `<N x iX>` (All ones in an element means true and all zeros means false). We need to perform a sign extension by copying the sign bit of the `i1` result until it reaches the size of `iX` (Program 3.4).

```
template <T> bitblock128_t simd<8>::add(bitblock128_t arg1, bitblock128_t arg2)
{
    return _mm_add_epi8(arg1, arg2);
}
```

Program 3.1: Implementation of `simd<8>::add` for X86 SSE2

```
template <> bitblock128_t simd<8>::add(bitblock128_t arg1, bitblock128_t arg2)
{
    return (bitblock128_t)vaddq_u8((uint8x16_t)(arg1), (uint8x16_t)(arg2));
}
```

Program 3.2: Implementation of `simd<8>::add` for ARM NEON

```
define <16 x i8> @simd_add_8(<16 x i8> %arg1, <16 x i8> %arg2) {
entry:
    %r = add <16 x i8> %arg1, %arg2
    ret <16 x i8> %r
}
```

Program 3.3: Implementation of `simd<8>::add` with LLVM IR

```
define <16 x i8> @simd_eq_8(<16 x i8> %arg1, <16 x i8> %arg2) {
entry:
    %r1 = icmp eq <16 x i8> %arg1, %arg2
    %r2 = sext <16 x i1> %r1 to <16 x i8>
    ret <16 x i8> %r2
}
```

Program 3.4: Implementation of `simd<8>::eq` with LLVM IR. `Sext` is the instruction for sign extension.

```
define <16 x i8> @simd_max_8(<16 x i8> %a, <16 x i8> %b) {
entry:
    %m = icmp sgt <16 x i8> %a, %b
    %r = select <16 x i1> %m, <16 x i8> %a, <16 x i8> %b
    ret <16 x i8> %r
}
```

Program 3.5: Implementation of `simd<8>::max` with LLVM IR. `Select` selects elements according to the first operand: $r_i = \begin{cases} a_i & \text{if } m_i = 1 \\ b_i & \text{otherwise} \end{cases}$.

- Operations that have no IR correspondence such as `simd::min` and `simd::max`. They can be emulated with a sequence of IR, e.g. `simd<8>::max` in Program 3.5.

For horizontal operations, IDISA also needs to maintain target-specific logic. For example, to implement `hsimd<16>::packh`, it uses unsigned saturation *packuswb* for X86 SSE2 and uses *vuzpq_u8* for NEON; for X86 SSE series after SSSE3, it uses the instruction *pshufb*. The author of the IDISA library needs to know these instruction sets very well. On the other hand, LLVM IR introduces a powerful instruction which can express most of the horizontal and expansion operations. It is the *shufflevector*.

```
<result> = shufflevector <n x <ty>> <v1>, <n x <ty>> <v2>, <m x i32> <mask>
; yields <m x <ty>>
```

The first two operands are vectors of the same type and their elements are numbered from left to right across the boundary. In the other word, the element indexes are $0 \dots n - 1$ for `v1` and $n \dots 2n - 1$ for `v2`. The `mask` is an array of constant integer indexes, which indicates the elements we want to extract to form the `result`. Either `v1` or `v2` can be "undefined" to do shuffle within one vector. *Shufflevector* is often used together with the *bitcast* operation. It converts between integer, vector and FP-values and changes the data type without moving or modifying the data, thus requiring the source and result type to have the same size in bits. With *shufflevector* and *bitcast*, we could write `hsimd<32>::packh` in Program 3.6. Figure 3.1 explains the indexes used in the shuffle mask.

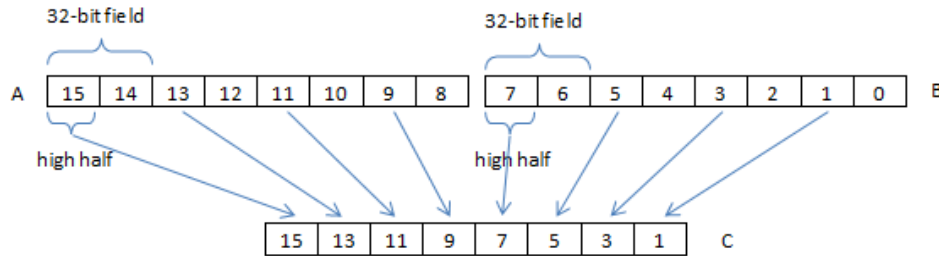


Figure 3.1: *Shufflevector* and `hsimd<32>::packh`. The vectors are bitcasted into *v8i16* and the indexes for the shuffle mask are drawn in the cell.

Program 3.6 can be easily generalized for packing high on any power-of-two field width. For other horizontal operations:

- Packing low: the same bitcast need to be done but *shufflevector* with a different mask. E.g. `hsimd<32>::packl` can be implemented with the mask 0, 2, 4, 6, 8, 10, 12, 14.
- Packing sign mask: it packs together all the sign bits from each field of the operand. This can be implemented with the less than comparison. E.g. `hsimd<32>::signmask(a)` is equivalent

```

define <8 x i16> @hsimd_packh_32(<4 x i32> %a, <4 x i32> %b) {
entry:
  %aa = bitcast <4 x i32> %a to <8 x i16>
  %bb = bitcast <4 x i32> %b to <8 x i16>
  %rr = shufflevector <8 x i16> %bb, <8 x i16> %aa, <8 x i32> <i32 1, i32 3,
    i32 5, i32 7, i32 9, i32 11, i32 13, i32 15>

  ret <8 x i16> %rr
}

```

Program 3.6: Shufflevector and `hsimd<32>::packh` in LLVM IR. Horizontal operations half the width of fields and that effect is reflected in the return value type.

to `icmp slt <4 x i32> %a, <4 x i32> <i32 0, i32 0, i32 0, i32 0>` which returns a `<4 x i1>` sign mask vector.

- Other operations that require coding a sequence of IR like `hsimd<32>::add_hl(a, b)`. They are rarely used in the Parabix application.

Shufflevector and bitcast could also cover IDISA expansion operations. We list the IR code for `esimd<16>::mergeh` in Program 3.7 and explain the indexes in Figure 3.2. The program is self-explanatory; any programmer who understands shufflevector can understand its behaviour easily.




Figure 3.2: Shufflevector and `esimd<16>::mergeh`. The indexes for the shuffle mask are drawn in the cell.

The rest of the expansion operations can be implemented as the following:

- Merge low: similar to merge high but with a different shuffle mask. E.g. `esimd<16>::mergel` uses the mask 0, 8, 1, 9, 2, 10, 3, 11.

```

define <4 x i32> @esimd_mergeh_16(<8 x i16> %a, <8 x i16> %b) {
entry:
    %rr = shufflevector <8 x i16> %b, <8 x i16> %a, <8 x i32> <i32 4, i32 12,
        i32 5, i32 13, i32 6, i32 14, i32 7, i32 15>

    %rr1 = bitcast <8 x i16> %rr to <4 x i32>
    ret <4 x i32> %rr1
}

```

Program 3.7: Shufflevector and `esimd<16>::mergeh` in LLVM IR. Expansion operations double the width of fields.

- Unary operations like sign extension and zero extension: LLVM has built-in instructions with the same name.

For field movement operations:

- Field extract or insert: LLVM IR offers two vector instructions `insertelement` and `extractelement` for them.
- Constant fill: it fills each field with an integer constant. In IR this can be coded with vector constants such as `<4 x i32> <i32 1, i32 10, i32 30, i32 99>`.
- Unary and binary movement: those operations move fields within one register or among two registers and they can be implemented with `shufflevector`.

The full register operations could be coded in large-size integers like `i128` and `i256`. You can add / multiply / shift it as a normal integer. In fact, all the integer instructions LLVM support can be applied to them thus enabling more complex operations that IDISA does not support.

To sum up, with support of all the five categories of IDISA operations, we are able to replace the IDISA library with pure IR implementation. However, there is still one question to answer: since LLVM has its own C++ compiler called Clang and Clang could compile the C++ IDISA library into IR, what is the difference between the Clang-generated IR and our hand-written IR library? There are at least three major difference:

1. Clang could not remove all the target-dependency from the C++ source. Not every IR instruction are target-independent. For example, IDISA function `hsimd<8>::packh` compiles to Program 3.8 and all the functions that start with `@llvm.x86.sse2` are only available on X86 SSE2. This is inherent to the use of direct compiler intrinsic in IDISA.

```

define <2 x i64> @hsimd_packh_8(<2 x i64> %a, <2 x i64> %b) #4 {
entry:
  %0 = bitcast <2 x i64> %a to <8 x i16>
  %1 = tail call <8 x i16> @llvm.x86.sse2.psrlw(<8 x i16> %0, i32 8) #1
  %2 = bitcast <2 x i64> %b to <8 x i16>
  %3 = tail call <8 x i16> @llvm.x86.sse2.psrlw(<8 x i16> %2, i32 8) #1
  %4 = tail call <16 x i8> @llvm.x86.sse2.packuswb.128(<8 x i16> %3, <8 x i16> %1) #1
  %5 = bitcast <16 x i8> %4 to <2 x i64>
  ret <2 x i64> %5
}

```

Program 3.8: Clang-generated IR for `hsimd<8>::packh` from compiling the IDISA function

2. Illegal operations are handled in different level. Compare the following IR implementation for `simd<4>::add`:

```
add <32 x i4> %a, %b
```

With the Clang-generated IR from the IDISA SSE2 file:

```

%and.i.i.i = and <2 x i64> %b, %m0
%0 = bitcast <2 x i64> %a to <16 x i8>
%1 = bitcast <2 x i64> %and.i.i.i to <16 x i8>
%add.i.i10.i = add <16 x i8> %0, %1
%2 = bitcast <16 x i8> %add.i.i10.i to <2 x i64>
%3 = bitcast <2 x i64> %b to <16 x i8>
%add.i.i.i = add <16 x i8> %0, %3
%4 = bitcast <16 x i8> %add.i.i.i to <2 x i64>
%and.i.i.i.i = and <2 x i64> %2, %m0
%and.i.i7.i.i = and <2 x i64> %4, %m1
%or.i.i.i.i = or <2 x i64> %and.i.i.i.i, %and.i.i7.i.i
ret <2 x i64> %or.i.i.i.i

```

Given that addition on `v32i4` is not supported by this target, the latter implements it with `v16i8` addition and a few logic operations in the front end level. Target information is required. And even if the latter code is migrated to a target that supports `v32i4` addition natively, it could not use that ability unless some fancy optimization could recognize the intention behind these 12 lines.

On the other hand, the former one-line code is not extended until the legalization phase in the back end. The target-specific details are thus left to the back end.

3. Since the illegal operation is not extended until the back end, more optimizations are available. The high-level intention of the IR instruction is better preserved. Still for `simd<4>::add`, if one of the operand is all zero, the front end could remove the single line of `add <32 x i4>` more easily than removing the 12 lines of code in the Clang-generated IR. It is useful for constant combination as well as other peephole optimizations because it simplifies the pattern recognition. We give an example of peephole optimization on the long integer shifting in Chapter 5.

This comparison explains our design goal: to replace the IDISA library with a high-level target-independent IR library. It is different from the IDISA approach fundamentally in the way that it tries not to instruct the compiler how to implement this operation, but rather tell what to implement.

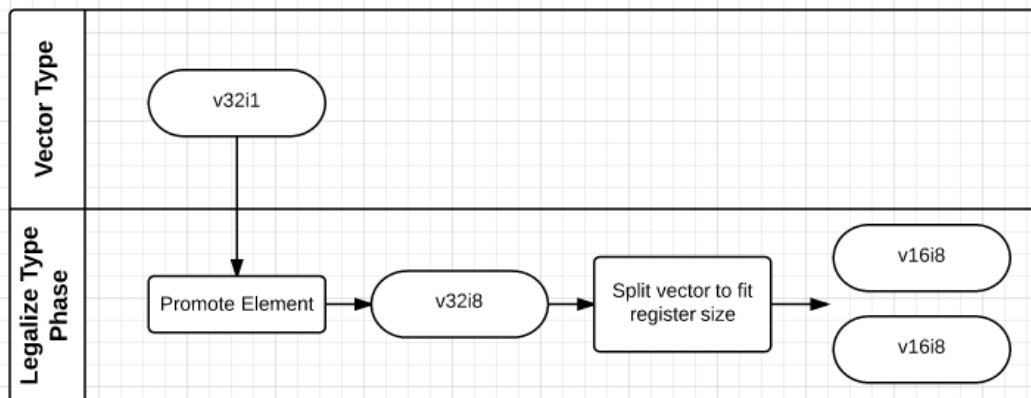
The next question would be that could LLVM compile the IR library to efficient machine code? Experiments on X86 tells no. Simple functions that directly respond to native instructions like `simd_add_8` in Program 3.3 can be compiled correctly, but for some more complexed ones like `shufflevector`, where no target so far has native support for it, poor machine code might be generated (e.g. pack high for 16-bit fields). Furthermore, LLVM does not have good support for vectors of small element, simple code like `add <128 x i1> %a, %b` would generate big amount of memory operations and eight additions on SIMD registers. To achieve a better code generation, we bring many of the strategies from IDISA to LLVM back end. The next two chapters would describe them in detail.

Chapter 4

Vector of $i2^k$

Parabix operation works on full range of vector types. For 128-bit SIMD register, Parabix supports $v128i1$, $v64i2$, $v32i4$, ..., $v1i128$, we call them as the vector of $i2^k$. Vector type $vXi8$, $vXi16$, ..., $vXi64$ is widely used for multimedia processing, digital signal processing and Parabix technology, they are well supported by the LLVM infrastructure, but the rest vector type with smaller element does not have perfect implementation. For instance, $vXi1$ is a natural view of many processor operations, like AND, OR, XOR; they are bitwise operations. However, $v32i1$, $v64i1$ and $v128i1$ are all illegal on current LLVM 3.4 back end for X86 architecture. After seeing a $v128i1$ vector, Type Legalize Phase would promote element type $i1$ to $i8$, and then split the vector to fit 128 bits register size; thus the incoming $v128i1$ turns into 8 $v16i8$ vectors. If we write AND on 2 $v128i1$ vectors, LLVM would produce 8 pairs of AND on $v16i8$ and also operations to truncate and concatenate back the $v128i1$ result; while we can simply bitcast $v128i1$ to any legal 128-bit vector like $v4i32$, do AND on them and bitcast the result back to $v128i1$. The performance penalty of type legalization is high in this example. Another type legalization example of $v32i1$ can be found in Figure 4.1.

LLVM applies the same promote element strategy to vectors of $i2$ and $i4$, which would lead to huge SelectionDAG generation and thus poor machine code. On the other hand, $i1$, $i2$ and $i4$ vectors are important to Parabix performance-critical operations, such as transposition and deletion; Parabix applications, such as DNA sequence (ATCG pairs) matching which can be encoded into $i2$ vectors most efficiently, requires a better support of small element vectors. The inductive doubling instruction set architecture (IDISA) which is the ideal model for Parabix needs a core set of functions on the $i2^k$ vectors as the first key element. All these reasons motivate us to find better implementation of $i1$, $i2$ and $i4$ vectors.

Figure 4.1: Type legalize process for $v32i1$ vector

4.1 Redefine Legality

In Chapter 2 we know that LLVM has three ways to legalize vector types: Scalarization, Vector Widening and Vector Element Promotion. None of these strategies would legalize small element vectors properly. Think about $v32i1$, it fits in the general 32-bit registers, and we can not benefit from extending or splitting the vector in wider or more registers, not to mention scalarizing it. It would be the best to store $v32i1$ vectors just in the general 32-bit register and properly handle the operations on them.

So we want to redefine the type legality inside LLVM. Instead of having direct hardware intrinsic on it, we define a vector type which has the same size in bits with one of the target's registers to be a legal vector type. The definition of the illegal operation remains the same. Under this definition, $v32i1$ is legal on any 32-bit platform, $v64i1$ is legal on any 64-bit platform and $v64i2$ is legal on any platform with 128-bit SIMD registers.

However, as more types are legal, we will need to handle more illegal operations. LLVM has the facility to "expand" an illegal operation, so that we do not need to implement every operation on the type. For example $v32i1$, we did not lower its shufflevector but we can still write shufflevector on $v32i1$ in IR, and LLVM could expand it into sequence of extracting and inserting vector elements. Of course the performance is not good at all. So the new question arises, what is the necessary operations set to fully support a legal type, that every possible IR statement with this type can be compiled into a native machine code?

This question is hard to answer. In practice, we implemented (1) common binary functions listed in Table 4.2; (2) basic vector operations like INSERT VECTOR ELT, EXTRACT VECTOR ELT and BUILD VECTOR. All the meaningful test IR files we wrote work properly under this operations set.

4.2 In-place Lowering Strategy

With the redefined legality, we provide the fourth way to legalize vector type: In-place Lowering. It is called "in-place" because we do not rearrange the bit value of the vector data, we would rather look at the same data with a different type. A trivial example would be the logical operations on $\langle 32 \times i1 \rangle$; we can simply bitcast $\langle 32 \times i1 \rangle$ to $i32$ and perform the same operation. Almost all the operations on $vXi1$ can be simulated with a few logic operations on iX (except the basic vector operations) as listed in Table 4.2. Figure 4.2 shows the overall process of lowering $v32i1$ addition.

In-place Lowering allows us to copy the vector between registers or shift the vector within the register boundary. But it is different from vector element promotion. Refer to the Figure 4.4, vector element promotion would require shifting different element with different offsets.




Figure 4.2: Comparison between LLVM default legalize process (left) and in-place lowering (right). The right marks $v32i1$ type legal and handles the operation ADD in the operation legalization phase. This will keep the data in the general registers without being promoted or expanded.

4.2.1 Lowering for $vXi2$

Vector type $vXi2$ has important role in the IDISA model and Parabix transposition and inverse transposition. Ideal Three-Stage Parallel Transposition[6] requires $hsimd<4>::packh$ and $hsimd<4>::packl$, which can be implemented with shufflevectors on $v64i2$. Shufflevectors of $v64i2$ are also required by Ideal Inverse Transposition, for $esimd<2>::mergeh$ and $esimd<2>::mergel$. Transposition is the first step of every parabix application[15] and it is the principle overhead for some application like regular expression matching[26]. So good code generation for $vXi2$ is important.

Lowering $vXi2$ is harder than $vXi1$, so we propose a systematic framework using logic and 1-bit shifting operations. Consider A, B as two $i2$ integers, $A = a_0a_1$ and $B = b_0b_1$, we can construct a truth table for every operation $C = OP(A, B)$. We then calculate the first bit and the second bit

Operation	Semantics
ADD	$c_i = a_i + b_i$
SUB	$c_i = a_i - b_i$
MUL	$c_i = a_i * b_i$
AND, OR, XOR	Common logic operations.
NE	Integer comparison between vectors. $c_i = 1$ if a_i is not equal to b_i .
EQ	$c_i = 1$ if a_i is equal to b_i
LT	$c_i = 1$ if $a_i < b_i$. a_i and b_i is viewed as signed integer
GT	$c_i = 1$ if $a_i > b_i$. a_i and b_i is viewed as signed integer
ULT	Same with LT, but numbers are viewed as unsigned integer
UGT	Same with GT, but numbers are viewed as unsigned integer
SHL	$c_i = a_i << b_i$. Element wise shift left
SRL	$c_i = a_i >> b_i$. Element wise logic shift right
SRA	$c_i = a_i >> b_i$. Element wise arithmetic shift right

Table 4.1: Supported operations and its semantics. A, B is the operands, C is the result. a_i, b_i, c_i is the i_{th} element.

Operation on $vXi1$	iX equivalence
ADD(A, B)	XOR(A', B')
SUB(A, B)	XOR(A', B')
MUL(A, B)	AND(A', B')
AND(A, B)	AND(A', B')
OR(A, B)	OR(A', B')
XOR(A, B)	XOR(A', B')
NE(A, B)	XOR(A', B')
EQ(A, B)	NOT(XOR(A', B'))
LT(A, B), UGT(A, B)	AND(A', NOT(B'))
GT(A, B), ULT(A, B)	AND(B', NOT(A'))
SHL(A, B), SRL(A, B)	AND(A', NOT(B'))
SRA(A, B)	A'

Table 4.2: Legalize operations on $vXi1$ with iX equivalence. A, B are $vXi1$ vectors, A', B' are iX bitcasted from $vXi1$. For $v128i1$, we use $v2i64$ instead of $i128$ since LLVM supports the former better.

of C separately with the logic combinations of a_0, a_1, b_0, b_1 and turn this into *Circuit Minimization Problem*: find minimized Boolean functions for c_0 and c_1 . We use Quine-McCluskey algorithm[19] to solve it; an example can be found in Table 4.3.

A	B	C
00	00	00
00	01	01
00	10	10
	...	
11	11	10

$$c_0 = (a_0 \oplus b_0) \oplus (a_1 \wedge b_1)$$

$$c_1 = a_1 \oplus b_1$$

Table 4.3: Truth table of ADD on 2-bit integers and the minimized Boolean functions for C .

Once we get the minimized Boolean functions, we can apply it onto the whole $vXi2$ vector. To do this, we need $IFH1(\text{Mask}, A, B)$ which selects bits from vector A and B according to the Mask . If the i_{th} bit of Mask is 1, A_i is selected, otherwise B_i is selected. $IFH1(\text{Mask}, A, B)$ simply equals to $(\text{Mask} \wedge A) \vee (\neg \text{Mask} \wedge B)$. With $IFH1$, if we have calculated the all high bits (c_0 for all the element) and low bits (c_1 for all the element), we can combine them with special HiMask , which equals to $101010 \dots 10$, 128 bits long in binary. To calculate all the high bits of each $i2$ element, we bitcast A, B into full register type (e.g. $v32i1$ to $i32$, $v64i2$ to $i128$ or $v2i64$) and then do the following substitution on the minimized Boolean functions:

- For a_0 and b_0 , replace it with A and B .
- For a_1 and b_1 , replace it with $A \ll 1$ and $B \ll 1$.
- Keep all the logic operations.

So $c_0 = (a_0 \oplus b_0) \oplus (a_1 \wedge b_1)$ becomes $(A \oplus B) \oplus ((A \ll 1) \wedge (B \ll 1))$, which simplifies to $(A \oplus B) \oplus ((A \wedge B) \ll 1)$. We use shifting to move every a_1 and b_1 in place. For all the lower bits of each $i2$ element, the rules are similar:

- For a_1 and b_1 , replace it with A and B .
- For a_0 and b_0 , replace it with $A \gg 1$ and $B \gg 1$.
- Keep all the logic operations.

Program 4.1 is the actual custom code to lower $v64i2$ addition. One thing to mention here is that we deploy a template system to automatically generate custom lowering code and the corresponding testing code. We would describe the template system later in Chapter 5.

```

static SDValue GENLowerADD(SDValue Op, SelectionDAG &DAG) {
    MVT VT = Op.getSimpleValueType();
    MVT FullVT = getFullRegisterType(VT);
    SDNodeTreeBuilder b(Op, &DAG);

    if (VT == MVT::v64i2) {
        SDValue A = b.BITCAST(Op.getOperand(0), FullVT);
        SDValue B = b.BITCAST(Op.getOperand(1), FullVT);

        return b.IFH1(/* 10101010...10, totally 128 bits */
            b.HiMask(128, 2),
            /* C0 = (A0 ^ B0) ^ (A1 & B1) */
            b.XOR(b.XOR(A, B), b.SHL<1>(b.AND(A, B))),
            /* C1 = (A1 ^ B1) */
            b.XOR(A, B));
    }

    llvm_unreachable("GENLower of add is misused.");
    return SDValue();
}

```

Program 4.1: The function generated to lower ADD on $v64i2$.

4.2.2 Inductive Doubling Principle For $i4$ Vector

Now we have better code generation for $vXi1$ and $vXi2$, $vXi4$ vectors are our next optimization target. Shufflevectors of $vXi4$ are used in `hsimd<8>::packh`, `hsimd<8>::packl` and `esimd<4>::mergeh`, which are required by Ideal Three-Stage Transposition / Inverse Transposition; $vXi4$ is also the critical part of the IDISA model. But unfortunately, the strategies discussed above cannot be applied to $vXi4$ efficiently.

Circuit Minimization Problem is NP-hard[8, 20]. For $vXi4$, we would have 4 Boolean functions of 8 variables: $c_i = f_i(a_0, a_1, a_2, a_3, b_0, b_1, b_2, b_3)$, $i \in \{0, 1, 2, 3\}$, and it is known that most Boolean functions on n variables have circuit complexity at least $2^n/n$ [20] and we need 1-bit, 2-bit, 3-bit shifting on A , B . So the framework on $vXi2$ could not generate efficient code for us at this time. Instead, we introduce *Inductive Doubling Principle* [15] and we will show that this general principle can be applied for $vXi4$ and even wider vector element type, e.g. multiplication on $v16i8$, to get better performance.

We use $v32i4$ as an example to illustrate Inductive Doubling Principle. To legalize $v32i4$, LLVM would promote this type into $v32i8$, widen every element to $i8$ and shift every element except the first one. Figure 4.4 shows an example of widening $v8i4$ into $v8i8$, we can see unnecessary movement of vector element during widening. On a platform with 128 bits SIMD register, $v32i8$ will further




Figure 4.3: To add 2 $v32i4$ vectors, a and b , we bitcast them into $v16i8$ vectors. The lower 4 bits of $A_0 + B_0$ gives us c_1 . We then mask out a_1 and b_1 (set them to zero), do add again, and the higher 4 bits of the sum is c_0 .




Figure 4.4: LLVM default type legalization of $v8i4$ to $v8i8$. a_0 to a_7 are $i4$ elements and they are shifted with different offsets during the element type promotion.

be divided into two $v16i8$ and take 2 registers to hold, while the original type $v32i4$ has 128 bits in size and should be able to reside in only 1 register. Inductive Doubling Principle could achieve the latter for us. It would bitcast the vector in-place, view the same register as $v16i8$ type and emulate $i4$ operations with $i8$; e.g. in Figure 4.3, to get add $\langle 32 \times i4 \rangle \%a, \%b$, we calculate c_0, c_2, \dots, c_{30} (high 4 bits in each $i8$ element) and c_1, c_3, \dots, c_{31} (low 4 bits in each $i8$ element) separately with 2 $v16i8$ additions:

$$C = IFH1(HiMask_8, A \wedge HiMask_8 + B \wedge HiMask_8, A + B) \quad (4.1)$$

$$HiMask_8 = (1111000011110000 \dots 11110000)_2 \quad (4.2)$$

So we can emulate SIMD operations on iX vectors with $iX/2$ or $i2X$ vector operations. We implemented all the operations on $vXi4$ with this principle and the algorithm is listed in Table 4.4. One thing needs to explain is SETCC, which is the internal representation of integer comparison

in LLVM. It has a third operand to determine comparison type, such as SETEQ (equal), SETLT (signed less than), and SETUGE (unsigned greater or equal to). The third operand preserves in our algorithm.

$$C = IFH1(HiMask_8, HiBits, LowBits)$$

Operation	$HiBits$	$LowBits$
$v32i4$	All operation is on $v16i8$	
MUL	$MUL(A \gg 4, B \gg 4) \ll 4$	Default
SHL	$SHL(A \wedge HiMask_8, B \gg 4)$	$SHL(A, B \wedge LowMask_8)$
SRL	$SRL(A, B \gg 4)$	$SRL(A \wedge LowMask_8, B \wedge LowMask_8)$
SRA	$SRA(A, B \gg 4)$	$SRA(A \ll 4, (B \wedge LowMask_8)) \gg 4$
SETCC	Default	$SETCC(A \ll 4, B \ll 4)$
Default OP	$OP(A \wedge HiMask_8, B \wedge HiMask_8)$	$OP(A, B)$

In the table:

$A \gg 4$: logic shift right every $i8$ element by 4 bits

$A \ll 4$: shift left of every $i8$ element by 4 bits

$HiMask_8 = (11110000 \dots 11110000)_2$

$LowMask_8 = (00001111 \dots 00001111)_2$

Table 4.4: Algorithm to lower $v32i4$ operations. The legalization input is $c = OP(a, b)$, where a, b, c are $v32i4$ vectors. A, B, C is the bitcasted results from a, b, c and they are all $v16i8$ type.

Furthermore, this method is applicable to vectors of wider element type. Multiplication on $v16i8$, for example, generates poor code on LLVM 3.4 (Program 4.2): the vectors are finally scalarized and 16 multiplications on $i8$ elements are generated. With in-place promotion, we bitcast the operands into $v8i16$ and generates 2 SIMD multiplications (*pmullw*) instead.

However, the algorithm in Table 4.4 cannot guarantee the best performance. Addition on $v32i4$ requires 2 $v16i8$ additions, but we can actually implement it with one. Look back to Figure 4.3, we need to mask out a_1, b_1 and do add again, because $a_1 + b_1$ may produce carry bit to the high 4 bits. If we mask out only the high bit of a_1 and b_1 , still we will not produce carry and we can calculate c_0 and c_1 together in one $v16i8$ addition. All we need to solve is how to put the high bit back. The following equations describe the 1-add algorithm:

$$m = (10001000 \dots 1000)_2 \quad (4.3)$$

$$A_h = m \wedge A \quad (4.4)$$

$$B_h = m \wedge B \quad (4.5)$$

$$z = (A \wedge \neg A_h) + (B \wedge \neg B_h) \quad (4.6)$$

$$r = z \oplus A_h \oplus B_h \quad (4.7)$$

Equation (4.6) uses only one $v16i8$ addition and equation (4.7) put the high bit back. Our vector legalization framework is flexible enough that we can choose to legalize $v32i4$ addition with 1-add

```

define <16 x i8> @mult_8(<16 x i8> %a, <16 x i8> %b) {
entry:
    %c = mul <16 x i8> %a, %b
    ret <16 x i8> %c
}

# LLVM 3.4 default:
pextrb $1, %xmm0, %eax
pextrb $1, %xmm1, %ecx

mulb    %cl
movzbl  %al, %ecx
pextrb  $0, %xmm0, %eax
pextrb  $0, %xmm1, %edx

mulb    %dl
movzbl  %al, %eax
movd    %eax, %xmm2
pinsrb  $1, %ecx, %xmm2
pextrb  $2, %xmm0, %eax
pextrb  $2, %xmm1, %ecx

mulb    %cl
movzbl  %al, %eax
pinsrb  $2, %eax, %xmm2
pextrb  $3, %xmm0, %eax
pextrb  $3, %xmm1, %ecx

mulb    %cl
movzbl  %al, %eax
pinsrb  $3, %eax, %xmm2
pextrb  $4, %xmm0, %eax
pextrb  $4, %xmm1, %ecx
...
...
(16 mulb blocks in total)

# Inductive doubling result:
movdqa  %xmm0, %xmm2
pmullw  %xmm1, %xmm2
movdqa  .LCPI0_0(%rip), %xmm3
movdqa  %xmm3, %xmm4
pandn   %xmm2, %xmm4
psrlw   $8, %xmm1
psrlw   $8, %xmm0
pmullw  %xmm1, %xmm0
psllw   $8, %xmm0
pand    %xmm3, %xmm0
por     %xmm4, %xmm0
retq

```

Program 4.2: Inductive doubling principle on $v16i8$ multiplication. LLVM 3.4 generate poor machine code, which will *pextrb* every $i8$ field and multiply them with *mulb*. We simplify it through 2 *pmullw*, which is the multiplication on $v8i16$.

algorithm while keeping the rest $v32i4$ operations under general in-place promotion strategy. We will discuss our framework implementation in Chapter 5.

4.3 LLVM Vector Operation of $i2^k$

In addition to the binary operations listed in Table 4.1, LLVM provides convenient vector operations like *insertelement*, *extractelement* and *shufflevector*, internally, they are DAG node INSERT VECTOR ELT, EXTRACT VECTOR ELT and VECTOR SHUFFLE. Another important internal node is BUILD VECTOR. In this section, we will discuss how to custom lower these nodes on $i2^k$ vectors.

BUILD VECTOR takes an array of scalars as input and output a vector with these scalars as elements. Take $v64i2$ vector on X86 SSE2 architecture for an example; ideally, the input would provide an array of 64 $i2$ scalars and BUILD VECTOR assembles them into a $v64i2$ vector. More specifically, since $i2$ is illegal on all X86 architecture, the legal input is actually 64 $i8$ scalars. The naive approach would be creating an "empty" $v64i2$ vector, truncating every $i8$ into $i2$ and inserting it into the proper location of the "empty" vector. We propose a better approach by rearranging the index.

Let us denote the input array as a_0, a_1, \dots, a_{63} , a_i is all $i8$. We rearrange them according to Table 4.5 and build 4 $v16i8$ vectors V_1, V_2, V_3, V_4 . The final build result is:

$$V = V_1 \vee (V_2 << 2) \vee (V_3 << 4) \vee (V_4 << 6) \quad (4.8)$$

a_{60}	\dots	a_{12}	a_8	a_4	a_0	V_1
a_{61}	\dots	a_{13}	a_9	a_5	a_1	V_2
a_{62}	\dots	a_{14}	a_{10}	a_6	a_2	V_3
a_{63}	\dots	a_{15}	a_{11}	a_7	a_3	V_4

Table 4.5: Rearranging index for BUILD VECTOR on $v64i2$

SIMD OR and SHL are used in this formula, thus improving the performance by parallel computing. Rearranging index approach can be easily generalized to fit BUILD VECTOR of $v128i1$ and $v32i4$.

EXTRACT VECTOR ELT takes 2 operands, a vector V and an index i . It returns the i_{th} element of V . The semantics would not allow much parallelism in the implementation. On X86 architecture, there are built-in intrinsic to extract vector element, such as *pextrb* ($i8$), *pextrw* ($i16$), *pextrd* ($i32$) and *pextrq* ($i64$); for smaller element type, we could extract the wider integer that contains it, shift the small element to the lowest bits and truncate. Following algorithm gives an example of extracting the i_{th} element from the $v64i2$ vector V .

- Bitcast V to $v4i32$ V' and extract the proper $i32$ E. Since every $i32$ contains 16 $i2$ elements, the

index of E is $\lfloor i/16 \rfloor$.

$$V' = \text{bitcast } \langle 64 \times i_2 \rangle V \text{ to } \langle 4 \times i_{32} \rangle$$

$$E = \text{extract element } V', \lfloor i/16 \rfloor$$

- Shift right E , to put the element we want in the lowest bits.

$$E' = E \gg (2 \times (i \bmod 16))$$

- Truncate the high bits of E' to get the result.

$$R = \text{truncate } i_{32} E' \text{ to } i_2$$

The choice of $v4i_{32}$ does not make a difference, we can use any of the wider element vector type mentioned above. On the X86 architecture, the support of extraction on $v8i_{16}$ starts at SSE2, while others start at SSE4.1, so we choose $v8i_{16}$ extraction in our code to target broader range of machines.

INSERT VECTOR ELT is similar, it takes 3 operands, a vector V , an index i and an element e . It inserts e into the i_{th} element of V and returns the new vector. Same as EXTRACT VECTOR ELT, X86 SSE2 supports $v8i_{16}$ insertion ($pinsrw$), SSE4.1 supports $v16i_8$ ($pinsrb$), $v4i_{32}$ ($pinsrd$) and $v2i_{64}$ ($pinsrq$); for smaller element type, we could extract the wider integer that contains the element, modify the integer and insert it back. Following algorithm gives an example of inserting e into the i_{th} element of the $v64i_2$ vector V .

- Bitcast V to $v4i_{32}$ V' and extract the proper i_{32} E .

$$V' = \text{bitcast } \langle 64 \times i_2 \rangle V \text{ to } \langle 4 \times i_{32} \rangle$$

$$E = \text{extract element } V', \lfloor i/16 \rfloor$$

- Truncate e and shift it to the correct position.

$$e' = \text{zero extend } (e \wedge (11)_2) \text{ to } i_{32}$$

$$f = e' \ll (2 \times (i \bmod 16))$$

- Mask out old content in E , put in the new element.

$$m = (11)_2 \ll (2 \times (i \bmod 16))$$

$$E' = (E \wedge \neg m) \vee f$$

- Insert back E' to generate the new vector R .

$$R = \text{insert element } V', E', \lfloor i/16 \rfloor$$

We have discussed VECTOR SHUFFLE in Chapter 3. We did not develop a general lowering strategy for the small element VECTOR SHUFFLE. In stead, we focused more on special cases that matter to Parabix critical operations, we optimized those cases to match performance of the hand-written library.

4.4 Long Stream Addition

Parabix technology has the concept of adding 2 unbounded streams and of course this needs to be translated into an block-at-a-time implementation[26]. One important operation is unsigned addition of 2 SIMD registers with carry-in and carry-out bit e.g. `add i128 %a, %b` or `add i256 %a, %b` with *i1* carry-in bit `c_in` and generates *i1* carry-out bit `c_out`. Dr Cameron developed a general model using SIMD methods for efficient long-stream addition up to 4096 bits in [26].

In this section, we will replace the internal logic of wide integer addition (*i128*, *i256* etc.) of LLVM with the Parabix long-stream addition. Same with Dr Cameron's work in [26], we assume the following SIMD operations on *i64* vectors legal on the target:

- `add <N x i64> X, Y`, where $N = \text{RegisterSize}/64$. SIMD addition on each corresponding element of the *i64* vectors, no carry bits could cross the element boundary.
- `icmp eq <N x i64> X, -1`: compare each element of X with the all-one constant, returning an `<N x i1>` result.
- `signmask <N x i64> X`: collect all the sign bit of *i64* elements into a compressed `<N x i1>` vector. From the LLVM speculation, this operation is equivalent to `icmp lt <N x i64> X, 0`, which is the signed less-than comparison of each *i64* element with 0. In the real implementation we use target-specific operations for speed, e.g. `movmsk_pd` for SSE2 and `movmsk_pd.256` for AVX.
- Normal bitwise logic operations on `<N x i1>` vectors. For small N, native support may not exist, so we bitcast `<N x i1>` to *iN* and then zero extend it to *i32*. This conversion could also help with the 1-bit shift we use later.
- `zext <N x i1> m to <N x i64>`: this corresponds to `simd<64>::spread(X)` in [26], which would distribute the N bits of the mask, one bit each to the lower end of the N *i64* elements.

We then present the long stream addition of 2 $N \times 64$ bit values X and Y with these operations as the following.

1. Get the vector sums of X and Y.

$$R = \text{add } \langle N \times i64 \rangle X, Y$$

2. Get sign masks of X , Y and R .

$$x = \text{signmask } \langle N \times i64 \rangle X$$

$$y = \text{signmask } \langle N \times i64 \rangle Y$$

$$r = \text{signmask } \langle N \times i64 \rangle R$$

3. Compute the carry mask c , bubble mask b and the increment mask i .

$$c = (x \wedge y) \vee ((x \vee y) \wedge \neg r)$$

$$b = \text{icmp eq } \langle N \times i64 \rangle R, -1$$

$$i = \text{MatchStar}(c*2+c_in, b)$$

`MatchStar` is a key Parabix operation which is developed for regular expression matching:

$$\text{MatchStar}(M, C) = (((M \wedge C) + C) \oplus C) | M$$

4. Compute the final result Z and carry-out bit c_out .

$$S = \text{zext } \langle N \times i1 \rangle i \text{ to } \langle N \times i64 \rangle$$

$$Z = \text{add } \langle N \times i64 \rangle R, S$$

$$c_out = i \gg N$$

One note here for the mask type: c and i are literally all $\langle N \times i1 \rangle$ vectors, but we actually bitcast and zero extend them into $i32$. This is useful in the formula $c*2+c_in$, `MatchStar` and $i \gg N$; in fact, after we shift left c by $c*2$, we already have an $N+1$ bit integer which will not fit in $\langle N \times i1 \rangle$ vector. The same is true for i ; so when we write $\text{zext } \langle N \times i1 \rangle i \text{ to } \langle N \times i64 \rangle$, there is an implicit truncating to get the lower N bits of i , but when we shift right i by $i \gg N$, we do not do such truncation.

LLVM internally implement long integer addition with a sequence of `ADDC` and `ADDE`, which is just chained 64-bit additions (or 32-bit additions on 32-bit target). We replace that with the long stream addition model thus improving the performance by parallel computing. As the hardware evolves, wider SIMD registers would be introduced, like 512-bit register in Intel AVX512, our general implementation could easily adopt this change in hardware and add two $i512$ in constant time.

During our implementation, we found there was no intrinsic in IR for addition with carry-in and carry-out bit, there was only one intrinsic `uadd.with.overflow` for addition with carry-out bit. To realize unbounded stream addition, the ability to take carry-in bit is necessary, otherwise we would end up with two `uadd.with.overflow` to include the carry-in bit. So we introduced a new intrinsic `uadd.with.overflow.carryin` and backed it with the long stream addition algorithm.

Chapter 5

Implementation

In this Chapter, we describe our realization of Parabix technology inside LLVM facility. LLVM is a well-structured open source compiler tool chain which is under rapid development. So during our implementation, we tried our best to follow its design principle while keeping our code modularized and isolated to be able to easily integrate with new versions of LLVM. Our goal of code design is to:

1. Use general strategies across different types and operations to reduce repeated logic.
2. Minimize code injection in the existing source and put Parabix logic in the separate module.
3. Put our code in auto-generated, thorough test.

Most of our code sits in LLVM Target-Independent Code Generator[4]. From Chapter 4, we know that current type legalization process of LLVM have big performance penalty for small element vectors, so our approach mark $i1$, $i2$ and $i4$ vector legal type first, and then handle them in Legalize Operation Phase. For convenience, we name this set of vector types *Parabix Vector*.

We walk through the following steps to mark a type legal on a certain target:

- **Add new register class in target description file.** LLVM uses TableGen (.td files) to describe target information which allows the use of domain-specific abstractions to reduce repetition [4]. Registers are grouped into register classes which further tie to a set of types. We introduced GR32X for 32-bit general register like EAX EBX for $v32i1$, GR64X for 64-bit general register like RAX RBX for $v64i1$, VR128PX for 128-bit vector register like XMM0 to XMM15 for $v128i1$, $v64i2$, $v32i4$ Types within the same register class can be bitcasted from one to the other, since they can actually reside in the same register.
- **Set calling convention.** They are two kinds of calling convention to set: return value and argument calling convention. For example, we instruct LLVM to assign $v64i2$ type return value

to XMM0 to XMM3 registers, assign $v64i2$ argument type to XMM0 to XMM7 registers if we have SSE2 or to 16-byte stack slots otherwise.

Now Legalize Type Phase recognizes our $i1$, $i2$, $i4$ vectors as legal and pass them onto Legalize Operation Phase. We have two major methods to handle $i2^k$ vectors here: *Custom Lowering* and *DAG Combining*.

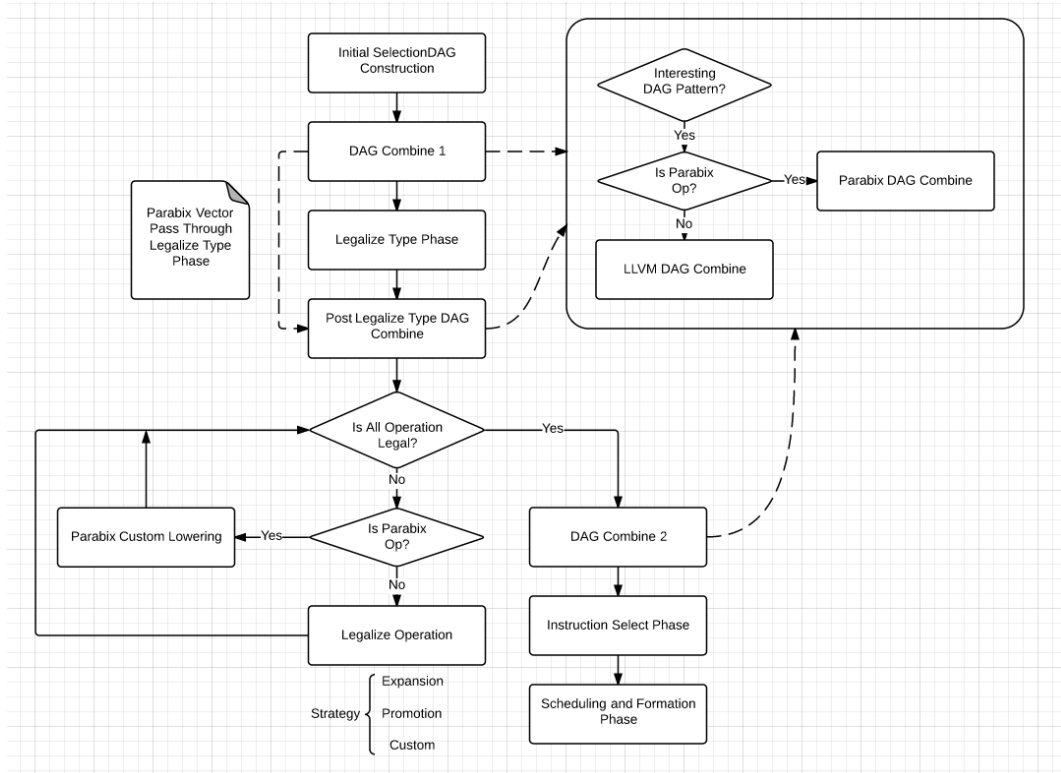


Figure 5.1: Overview of the modified instruction selection process. Logic for Parabix vectors are hooked into two places: the Legalize Operation Phase and DAG Combine Phases. Parabix Custom Lowering and Parabix DAG Combine are both modularized and separated.

Figure 5.1 gives an overview of our implementation. Custom Lowering resides in the operation legalizer and LLVM does it through iteration, which allows the legalizer to introduce new illegal operations within each iteration. Every time the legalizer finds an illegal operation, it will check if that is a Parabix operation and if so redirect to the Parabix Custom Lowering module. DAG Combining is similar but we have multiple combine timings available and it is designed mainly for cleaning up the messy output of the legalizers.

In the following sections, we will discuss some custom lowering strategies and how they are organized to fit our design goal; then we give some examples of the Parabix DAG Combiner which are usually special cases for a certain operation; finally we will show how we use templates to

generate code and test cases for the sake of DRY (don't repeat yourself).

5.1 Standard Method For Custom Lowering

5.1.1 Custom Lowering Strategies

After the Legalize Type Phase, one shall not generate illegal types again. This means all the phases after type legalization are target-specific. But in practice, almost all the targets support $i8$, $i32$ and $i64$, so there are still general strategies we can apply across targets. For different types like $v32i1$ and $v128i1$, general strategies also exist to lower both of them. We define three legalize actions as the following:

1. Bitcast to full register and replace operation code. This is useful for all $i1$ vectors, we need to specify the new operation code when defining the action, e.g. XOR for ADD on $v32i1$.
2. In-place promotion. Automatically apply $i2X$ vector operations on iX vector following the Inductive Doubling Principle.
3. Custom. Same concept with LLVM Custom Lowering, manually replace an illegal DAG node with a sequence of new DAG nodes. They can be illegal nodes, but they cannot introduce illegal types. All $i2$ vectors are lowered here, also the 1-add version of the $v32i4$ addition.

5.1.2 DAG Combiner

DAG Combiner is the supplement to custom lowering facility and it often focuses on special cases, e.g. one operation and a subset of possible operands. We give a few examples of Parabix DAG Combiner here.

The first example is shufflevector for packh("pack high") and packl("pack low"). LLVM 3.4 does not generate the best assembly code for $v16i8$ packing, it generates a sequence of *pextrw* and *pinsrw*. So we create the following DAG Combiner:

- **Pattern:** shufflevector on $v16i8$ with mask = 0, 2, 4, ..., 30 (packl) or mask = 1, 3, 5, ..., 31 (packh).
- **Combine Result:** one PACKUS node, which would unsigned saturate two $v8i16$ into $v8i8$ vectors and concatenate them into one $v16i8$.

Furthermore, since $v128i1$, $v64i2$ and $v32i4$ are legal vectors now, we have the chance to optimize shufflevector on them too. Still use packh/packl as an example, we can utilize the PEXT node introduced by the Intel Haswell Architecture, BMI2. PEXT is an useful instruction for bit manipulation on $i32$ and $i64$. Given the $i8$ variable $A = abcdefgh$, $Mask = (10101010)_2$, $PEXT(A, Mask)$

```

define <2 x i64> @packh_2(<2 x i64> A, <2 x i64> B) {
entry:
    ; extract lower 64 bits (A0) and higher 64 bits (A1)
    A0 = extractelement <2 x i64> A, i32 0
    A1 = extractelement <2 x i64> A, i32 1

    Mask = 0xAAAAAAAAAAAAAAAA ; 1010...1010 in binary
    P0 = PEXT(A0, Mask) | (PEXT(A1, Mask) << 32)

    ; same for B
    B0 = extractelement <2 x i64> B, i32 0
    B1 = extractelement <2 x i64> B, i32 1
    P1 = PEXT(B0, Mask) | (PEXT(B1, Mask) << 32)

    ret <2 x i64> <i64 P0, i64 P1>
}

```

Program 5.1: Implementation of `hsimd<2>::packh` with PEXT.

would return $R = aceg$ (a,b,c,d,e,f,g are single bits). So PEXT would extract bits from A at the corresponding bit locations specified by Mask. With this in mind, we can implement `hsimd<2>::packh` as Program 5.1; for readability, it is in pseudo IR.

According to Program 5.1, we create the following DAG Combiner:

- **Pattern:** shufflevector on v_{128i1} , v_{64i2} or v_{32i4} with mask = 0, 2, 4, ..., $NumElt \times 2 - 2$ or mask = 1, 3, 5, ..., $NumElt \times 2 - 1$. $NumElt$ is the number of elements for each type e.g. $NumElt = 32$ for v_{32i4} .
- **Combine Result:** four PEXT nodes combined with OR and SHL.

To summarize, this kind of DAG Combiner provides a shortcut for the programmer to do ad hoc optimization and it can co-exist with a full custom lowering, like the relationship between immediate shifting and arbitrary shifting. Immediate shifting shifts all the vector elements with the same amount and we can have efficient realization for v_{32i4} with v_{4i32} shifts, while we apply In-place Promotion strategy for v_{32i4} arbitrary shifting in the Parabix custom lowering.

Apart from this, DAG Combiner can also optimize operations with illegal type in the phase DAG Combine 1, which is not possible in the custom lowering. But we cannot simply put all the Parabix Custom Lowering logic inside the DAG Combiner. First, it is against LLVM design; DAG Combiner is designed for cleaning up, either the initial code or the messy code generated by the legalization passes [4]. Second, it cannot utilize the legalization iteration; in custom lowering, general strategies may introduce new illegal operations and they are hard to avoid since "illegal" is a target-specific

concept; these illegal operations will be lowered in the next iteration and so on. The DAG Combiner, on the other hand, 1) Should not generate illegal operations in the phases after the Legalize Operation Phase. 2) Although it can also work in iteration, most of the lowering logic for common operations are not programmed in this module, we would end up with illegal non-Parabix operations.

5.2 Templated Implementation

During our implementation, we encountered many duplicated code, especially in the test cases; they are against software design principles and they are hard to maintain, sometimes even hard to write: a thorough test file for $i2^k$ vector could contain more than two thousand lines of code, most of which are in the same pattern. To keep DRY (don't repeat yourself) and save programmer time, we introduced Jinja template engine [3]. According to [7], Jinja belongs to the Engines Mixing Logic into Templates, it allows embedding logic or control-flow into template files. We use Jinja because:

- We can write all pieces of the content in one file, so it is easier to understand. Where in the Engines using Value Substitution, the driver code usually contains many tiny pieces of content, and the reader must read the driver code as well as the template file to understand the output. Examples can be found in Program 5.2.
- Like the standard Model-View-Controller structure in the web design, our driver code needs only provide abstract data (like operation names in IR and the corresponding C++ library calls), and how to present these data is not its responsibility. In the other word, we can have significant changes in the template without changing the driver.
- Jinja uses python and python is easy and quick to use.

5.2.1 Code Generation For $i2$ Vector

In Chapter 4, we legalized $i2$ vector operations with boolean functions. In our implementation, with the Quine-McCluskey solver, we got 11 sets of formula which reside in one compact data script and we wrote template files to generate 11 C++ functions for them. This approach allows us to:

- Collect all the critical formula together so that possible future updates are easy to deploy.
- Implementation details only reside in the template file, so we are able to change the code structure easily. For now we create one function for each formula, but it is possible that we plan to create one big switch statement and generate one case for each formula instead. This can be done with only a few lines of change in the template.

5.2.2 Test Code And IR Library Generation

To test the vector of $i2^k$, we need a pure IR library and a test driver. We could compile the IR library into one object file and link it with the driver, so that the driver could generate random test data, pass them to both the IR library and the reference library IDISA+ [18], and compare their results. The test system overview can be found in Figure 5.2. We use templates for both the IR library and the driver, some sample templates can be found in Program 5.2.

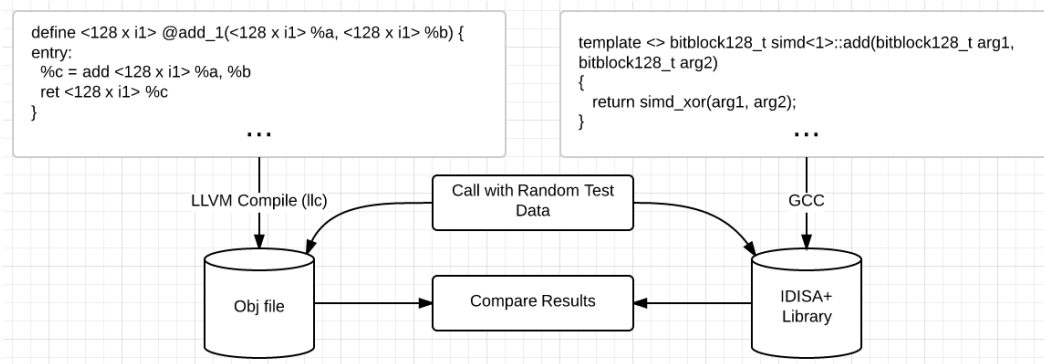


Figure 5.2: Test system overview. The pure IR library is first compiled into the native object file and then linked with the driver. The driver call functions from the both side to check correctness.

```

{% for name in FunctionNamesI4 %}
define <32 x i4> @{{name.c}}(<32 x i4> %a,
                           <32 x i4> %b)
{
entry:
    %c = {{ name.op }} <32 x i4> %a, %b
    {% if "icmp" in name.op %}
    %d = sext <32 x i1> %c to <32 x i4>
    ret <32 x i4> %d
    {% else %}
    ret <32 x i4> %c
    {% endif %}
}
{% endfor %}

```

```

define <32 x i4> @add_4(<32 x i4> %a,
                      <32 x i4> %b)
{
entry:
    %c = add <32 x i4> %a, %b
    ret <32 x i4> %c
}

```

```

define <32 x i4> @eq_4(<32 x i4> %a,
                     <32 x i4> %b)
{
entry:
    %c = icmp eq <32 x i4> %a, %b
    %d = sext <32 x i1> %c to <32 x i4>
    ret <32 x i4> %d
}

```

Program 5.2: Templates for the IR Library. On the top is the template, and two different output are listed below. We use embedded for loop and if statements.

Chapter 6

Performance Evaluation

In this chapter, we focus on the performance evaluation to show that our LLVM back end could not only match performance with the hand-written library, but also provide a better chance to optimize according to the specific target. We would first validate our vector of $i2^k$ approaches, and then present the performance of some critical Parabix operations via application-level profile.

6.1 Vector of $i2^k$ Performance

In Chapter 4, we present different approaches to lower $i1$, $i2$, $i4$ and some $i8$ operations within one SIMD register. In this section, we would validate our approaches by showing the improved run-time performance.

6.1.1 Methodology

Testing small pieces of critical code can be tricky, since the testing overhead can easily overwhelm the critical code and make the result meaningless. Agner Fog provides a test program which uses the Time Stamp Counter for clock cycles and Performance Monitor Counters for instruction count and other related events [9]. We pick the reciprocal throughput as our measurement and it is measured with a sequence of same instructions where subsequent instructions are independent of the previous ones. In Fog's instruction table, he noted that a typical length of the sequence is 100 instructions of the same type and this sequence should be repeated in a loop if a larger number of instructions is desired.

We did one simple experiment with SIMD XOR (*xorps*) to validate this program. Refer to Figure 6.1, we measured the performance of executing different number of XOR instructions; they are organized into one for loop and we have checked the assembly code to make sure the XOR

operations are not optimized away.




Figure 6.1: Test performance with XOR. The dotted line is instruction count and the other line is core CPU cycles.

From the figure, we can see the instruction count and CPU cycles grows linearly with the number of XOR instructions. So we can conclude that Fog's test program can be used to compare two pieces of critical code: the one with more measured CPU cycles is more complex and have more instructions. Note that from the figure, it seems the throughput of *xorps* is 4, which is different from Intel's document (3 in document). We found this may be related to the compiler optimization on the loop; when we flattened the loop we got the throughput around 2 to 3. In order to eliminate this undesired effect, we would flatten the test code in the following sections.

In the following sections, we would write micro benchmarks with Agner Fog's test program and compare reciprocal throughput between different implementation. Our test machine is X86 64-bit Ubuntu with Intel Haswell, and the detailed configuration can be found in Table 6.1. In order to inline pure IR functions (instead of a function call into one object file), we compile all the test code into LLVM bit code (binary form of LLVM IR) and then link / optimize them together. The default compile flag is to use Intel SSE2 instruction set and on 64-bit OS.

6.1.2 Performance Against IDISA

We compare our lowering on pure IR function with IDISA+ Library [18] which is written in C++. To test each operation, we generate a sequence of 500 such operations where none them has to wait for the previous one. 100 operations seems not long enough for a stable result. The performance

CPU Name	Intel(R) Core(TM) i5-4570 CPU
CPU MHz	3200
FPU	Yes
CPU(s) enabled	4 cores
L1 Cache	32 KB D + 32KB I
L2 Cache	256 KB
L3 Cache	6 MB
Memory	8GB

Table 6.1: Hardware Configuration

Operating System	Ubuntu (Linux X86-64)
Compiler	Clang 3.5-1ubuntu1, GCC 4.8.2
LLVM	LLVM 3.4
File System	Ext4

Table 6.2: Software Configuration

comparison is listed in Figure 6.2 and Figure 6.3. We can see for $i1$ and $i4$ vectors, IR library has the similar performance with IDISA but it performs better with $i2$ vectors, especially on integer comparison.

The underlying logic for both libraries is the same, but it is implemented in different level. For IDISA library, `simd<2>::ugt` got inline-extended immediately by the compiler front end and its semantics of integer comparison lost ever after, while in the IR library, for the whole life cycle before the instruction selection, `ugt_2` keeps its semantics and this may help the compiler to optimize. The expansion of `ugt_2` is delayed until the instruction selection phase, right before machine code generation. We checked that IDISA function `simd<2>::ugt` and IR function `ugt_2` (whose underlying code is just `icmp ugt <64 x i2> %a, %b`) generated different assembly code.

However, the delay in expansion is not always good. Take multiplication on the $i2$ vector for an example, we can see our IR library has slightly better total CPU cycles, but if we write our instructions sequence with a loop, IDISA library would win (Figure 6.4). Loop optimization should be responsible for this difference and we did observe some kind of hoisting in the assembly code. Early expansion in IDISA also provides more optimization opportunity to the compiler front end.

Further more, from the reciprocal throughput comparison (Figure 6.2), IR library loses a bit on $i1$ vectors but wins most of the cases in $i2$ and $i4$; it may relate to a better instruction selection. IDISA library is generated from a strategy pool based on the number of basic instructions which are treated equally with cost 1. But basic instructions actually have different throughput in the real hardware, and LLVM back end are aware of that, thus selecting better instructions.

6.1.3 Performance Against LLVM

We compare our lowering with native LLVM. LLVM could not handle $i2$, $i4$ vectors and handle $i1$ vectors slowly. Detailed performance data can be found in Table 6.3. We can see that our approach fills the gap of LLVM type system.

	$i1$	$i2$	$i4$	$i8$
add	302	X	X	1
sub	310	X	X	1
mult	X	X	X	10
eq	273	X	X	1
lt	X	X	X	1
gt	X	X	X	1
ult	349	X	X	1
ugt	290	X	X	1

Table 6.3: Performance against LLVM native support of $i2^k$ vectors. 'X' means compile error or compile too slowly (longer than 30s), the rest number means the ratio of CPU cycles speed up: add takes 302 times of cycles that our lowering needs. For $i8$, we apply inductive doubling strategy on the multiplication, which explains the 10 times speed up.

6.2 Parabix Critical Operations

In this section, we evaluate our work by replacing Parabix critical operations with the IR library. We first choose transposition and inverse transposition as two representative operations and measure performance in two Parabix applications: XML validator and UTF-8 to UTF-16 transcoder. Note that we did not rewrite the whole application with an IR library, part of the application is still IDISA but some critical operation is replaced. The default compile flag is to use Intel SSE2 instruction set and on 64-bit OS.

	dew.xml	jaw.xml	roads-2.gml	po.xml	soap.xml
xmlwf0	3.93	4.364	4.553	4.891	5.18
xmlwf0 on Haswell	3.929	4.363	4.554	4.876	5.178
xmlwf1	3.929	4.371	4.566	4.861	5.186
xmlwf1 on Haswell	3.566	3.978	4.163	4.451	4.787

Table 6.4: Performance comparison of XML validator (xmlwf), in a thousand CPU cycles per thousand byte. In the table, xmlwf0 is implemented with full IDISA library and xmlwf1 is a copy of xmlwf0 with the transposition replaced.

Table 6.4 shows the performance of the XML Validator. The only difference of xmlwf0 and xmlwf1 is their transposition code, and the one in xmlwf1 is written in pure IR with the byte-pack algorithm. We can see xmlwf0 and xmlwf1 share almost identical performance and it is not for free. LLVM 3.4

	dew.xml	jaw.xml	roads-2.gml	po.xml	soap.xml
U8u16_0	281.46	37.11	40.06	244.94	10.2
U8u16_0 Haswell	272.68	34.21	39.84	242.56	10.11
U8u16_1	284.17	36.71	41.65	255.57	10.6
U8u16_1 Haswell	267.14	34.64	38.53	237.66	9.98

Table 6.5: Performance comparison of UTF-8 UTF-16 transcoder, in a million CPU cycles. U8u16_0 is written in IDISA, U8u16_1 has the transposition and inverse transposition part replaced.

cannot handle packing on 16-bit field width very well and we custom lower the shufflevector and generate PACKUS instruction for X86.

Another interesting observation is, when we re-compiled the same code on the Intel Haswell platform, we got almost no improvement for xmlwf0, since the IDISA library linked in is written with SSE2 intrinsic and only SSE2 instructions can be generated; but we got a slightly better performance for xmlwf1, because the IR library is target-independent and LLVM back end knows other instruction sets like SSE3, SSE4 is available on this platform, it generates better code with them.

Similar performance data on UTF-8 to UTF-16 transcoder is listed in Table 6.5. U8u16_0 is written in IDISA and U8u16_1 has both the transposition and inverse transposition part replaced. We also tried to compile them on the full Haswell, which gave us similar performance benefit.

6.2.1 Ideal 3-Stage Transposition on the Intel Haswell

Intel Haswell architecture introduces PEXT operation which can be used for the ideal 3-stage transposition. We evaluated its performance in Table 6.6. We can see the performance drops with PEXT, but the major reason is that PEXT can only work on *i32* or *i64* integer for the current architecture, not the algorithm. As the hardware evolves, we may have PEXT on SIMD registers directly and we can expect a better performance in xmlwf2, may be better than both xmlwf0 and xmlwf1 since 3-stage transposition is proved to be optimal under the IDISA model [15]. Our approach provides a new chance to exploit future hardware benefit without changing the source code.

	dew.xml	jaw.xml	roads-2.gml	po.xml	soap.xml
xmlwf0 on Haswell	3.929	4.363	4.554	4.876	5.178
xmlwf1 on Haswell	3.566	3.978	4.163	4.451	4.787
xmlwf2 on Haswell	4.11	4.49	4.69	4.978	5.308

Table 6.6: Ideal 3-stage transposition on xmlwf2. Xmlwf1 uses byte-pack algorithm in IR, xmlwf0 uses the same algorithm in IDISA.

6.2.2 Long Stream Addition

We replaced the internal logic of big integer addition in Chapter 4 and introduced a new intrinsic: `uadd.with.overflow.carryin`. We would evaluate them in this section by first comparing the long-stream addition algorithm with LLVM's original implementation and then doing some application level profile for the new intrinsic.

We wrote micro benchmarks with Fog's test program and we put 200 independent additions on *i128* and *i256*. It was tricky to make the test program right, we generated random data for the operands and we carefully inserted the carry-out bit back to the return value so that our long stream addition logic would not be optimized away. In order to be consistent throughout the comparison, we used the same compiler flag for all the runs (`-mavx2` for gcc and `-mattr=+avx2,+bmi2` for LLVM tool chain). The result is listed in Table 6.7.

	Core CPU Cycles	Instructions
Long stream addition on <i>i128</i>	2416	6552
LLVM on <i>i128</i>	1455	4199
Long stream addition on <i>i256</i>	2656	6959
LLVM on <i>i256</i>	4234	9798

Table 6.7: Micro benchmarks for long stream addition against LLVM's original implementation.

Long stream addition does not perform well on *i128*, since there is only two sequential additions involved (1 *addq* and 1 *adcq*) and parallel computing would not save much but introduce complexity. However, we can see on *i256* long stream addition has better performance than the sequential one that generates 1 *addq* and 3 *adcq*. As the width of the operand doubles, the CPU cycles from LLVM increases to the rate of 2.91, while in the long stream addition, the rate is only 1.10. Our algorithm scales better when the width of SIMD registers grows. We could confidently predict that on the Intel AVX512, long stream addition on *i512* would out-perform the sequential one significantly.

We then show application level profile of 'icgrep' which is a tool for regular expression matching with bitwise data parallelism. The internal "add with carry" logic is replaced with one single intrinsic and the performance is plotted in Figure 6.5. The performance drops because that version of icgrep works with 128-bit SIMD registers and long stream addition does not work well on *i128*. We could expect better performance on wider SIMD registers.




Figure 6.2: Reciprocal instruction throughput against IDISA library. IR and IDISA share almost identical throughput.



Figure 6.3: Total CPU cycles against IDISA library; for *i1* and *i4* vectors, IR library has the similar performance with IDISA but it performs better with *i2* vectors, especially on integer comparison.




Figure 6.4: The same benchmark for *i2* vectors with the instruction in a loop. Code in Figure 6.3 can be seen as the flattened version of this figure. We find IDISA here wins in the multiplication on *i2*, while IR wins it in Figure 6.3. Loop optimization should be responsible for it.




Figure 6.5: Performance of icgrep with long stream addition. This version of icgrep is based on *i128* so long stream addition actually slows down the performance; but for wider SIMD register, it would get the same performance or even better. For different regular expressions, the portion of "add with carry" code can be different, which explains the difference of the relative cycles across the x axis.

Chapter 7

Conclusion

In this thesis, LLVM as a new back end is introduced to the Parabix technology and a target-independent IR library of critical Parabix operations has been developed. Introducing LLVM brings in mature inter-procedure optimization and just-in-time compilers and outsources the machine-level code generation from the Parabix framework.

A systematic support for the vector of $i2^k$ has also been developed to extend the LLVM code generator with the IDISA model and a new LLVM intrinsic is added to enable chained additions on unbounded bit streams, which can be used for a broad category of applications. In one specific target: Intel X86, efficient native code has been generated and the performance is as good as the well-tuned IDISA library. In some micro-benchmarks, it even achieves 300 times speedup over the unmodified LLVM. Performance improvement over different sub-targets (e.g. X86 SSE2 and AVX2) has been witnessed without any change in the IR library.

With the LLVM back end, new optimization passes for the Parabix can be developed. As one of the major reasons for its high performance, Parabix uses long sequence of bitwise logic and shift operations without any branch or loop statement. Specific optimization like new register allocation algorithm may benefit this style of programming much more than the general programs. A peephole optimizer may be added to LLVM to combine a sequence of operations into one compact SIMD intrinsic available on the particular target.

Parabix with LLVM has better chances to target at different platforms efficiently such as the SPARC servers from the Sun and the ARM mobile platform. Further extension of the LLVM code generator can be done in the future for these platforms.

Bibliography

- [1] LLVM Language Reference Manual . <http://llvm.org/docs/LangRef.html>. 11, 12
- [2] IDISA toolkit project. <http://parabix.costar.sfu.ca/wiki/IDISAProject>. 8
- [3] Jinja documentation. <http://jinja.pocoo.org/>. 38
- [4] The LLVM target-independent code generator. <http://llvm.org/docs/CodeGenerator.html>. 11, 12, 34, 37
- [5] LLVM's Analysis and Transform Passes. <http://llvm.org/docs/Passes.html>. 10
- [6] The parabix transposition. <http://parabix.costar.sfu.ca/wiki/ParabixTransform>. 23
- [7] Python templating. <https://wiki.python.org/moin/Templating>. 38
- [8] Quine McCluskey algorithm wikipedia. http://en.wikipedia.org/wiki/Quine_McCluskey_algorithm. 26
- [9] Test programs for measuring clock cycles and performance monitoring. <http://www.agner.org/optimize/>. 41
- [10] Yosi Ben Asher and Nadav Rotem. Hybrid type legalization for a sparse SIMD instruction set. *ACM Trans. Archit. Code Optim.*, 10(3):11:1–11:14, September 2008. 1, 12
- [11] Robert D Cameron. u8u16—a high-speed utf-8 to utf-16 transcoder using parallel bit streams. Technical report, Technical Report TR 2007-18, Simon Fraser University, Burnaby, BC, Canada, 2007. 1
- [12] Robert D Cameron. A case study in simd text processing with parallel bit streams: Utf-8 to utf-16 transcoding. In *Proceedings of the 13th ACM SIGPLAN Symposium on Principles and practice of parallel programming*, pages 91–98. ACM, 2008. 1, 9, 10
- [13] Robert D Cameron, Ehsan Amiri, Kenneth S Herdy, Dan Lin, Thomas C Shermer, and Fred P Popowich. Parallel scanning with bitstream addition: An xml case study. In *Euro-Par 2011 Parallel Processing*, pages 2–13. Springer, 2011. 6, 7
- [14] Robert D Cameron, Kenneth S Herdy, and Dan Lin. High performance xml parsing using parallel bit stream technology. In *Proceedings of the 2008 conference of the center for advanced studies on collaborative research: meeting of minds*, page 17. ACM, 2008. 1
- [15] Robert D. Cameron and Dan Lin. Architectural Support for SWAR Text Processing with Parallel Bit Streams: The Inductive Doubling Principle. *SIGPLAN Not.*, 44(3):337–348, March 2009. 1, 7, 9, 23, 26, 45

- [16] Ron Cytron, Jeanne Ferrante, Barry K Rosen, Mark N Wegman, and F Kenneth Zadeck. Efficiently computing static single assignment form and the control dependence graph. *ACM Transactions on Programming Languages and Systems (TOPLAS)*, 13(4):451–490, 1991. 10
- [17] Randall James Fisher. *General-purpose Simd Within a Register: Parallel Processing on Consumer Microprocessors*. PhD thesis, West Lafayette, IN, USA, 2003. AAI3108343. 1
- [18] Hua Huang. IDISA+: A portable model for high performance SIMD programming. Master's thesis, Simon Fraser University, December 2011. 2, 4, 5, 8, 39, 42
- [19] E.D. Johnson. *Quine-McCluskey: A Computerized Approach to Boolean Algebraic Optimization*. Massachusetts Institute of Technology, Department of Electrical Engineering and Computer Science, 1981. 25
- [20] Valentine Kabanets and Jin-Yi Cai. Circuit minimization problem. In *Proceedings of the thirty-second annual ACM symposium on Theory of computing*, pages 73–79. ACM, 2000. 26
- [21] Chris Lattner. LLVM: An Infrastructure for Multi-Stage Optimization. Master's thesis, Computer Science Dept., University of Illinois at Urbana-Champaign, Urbana, IL, Dec 2002. See <http://llvm.cs.uiuc.edu>. 2, 10
- [22] Chris Lattner and Vikram Adve. LLVM: A Compilation Framework for Lifelong Program Analysis & Transformation. In *Proceedings of the 2004 International Symposium on Code Generation and Optimization (CGO'04)*, Palo Alto, California, Mar 2004. 2, 10
- [23] Ruby B. Lee. Accelerating multimedia with enhanced microprocessors. *IEEE Micro*, 15(2):22–32, 1995. 1
- [24] Nigel Woodland Medforth. icxml: Accelerating xerces-c 3.1. 1 using the parabix framework. 2013. 1
- [25] Huy Nguyen and Lizy Kurian John. Exploiting simd parallelism in dsp and multimedia algorithms using the altivec technology. In *Proceedings of the 13th International Conference on Supercomputing, ICS '99*, pages 11–20, New York, NY, USA, 1999. ACM. 1
- [26] Arrvindh Shriraman Kenneth S. Herdy Dan Lin Benjamin R. Hull Meng Lin Robert D. Cameron, Thomas C. Shermer. Bitwise data parallelism in regular expression matching. 1, 6, 7, 10, 23, 32
- [27] David A Terei and Manuel MT Chakravarty. An llvm backend for ghc. In *ACM Sigplan Notices*, volume 45, pages 109–120. ACM, 2010. 2, 10